

# Certified Kubernetes Administrator

Drills

# RX-M Cloud Native Advisory, Consulting and Training

2

Copyright 2013-2019, RX-M LLC

## ▪ Microservice Oriented

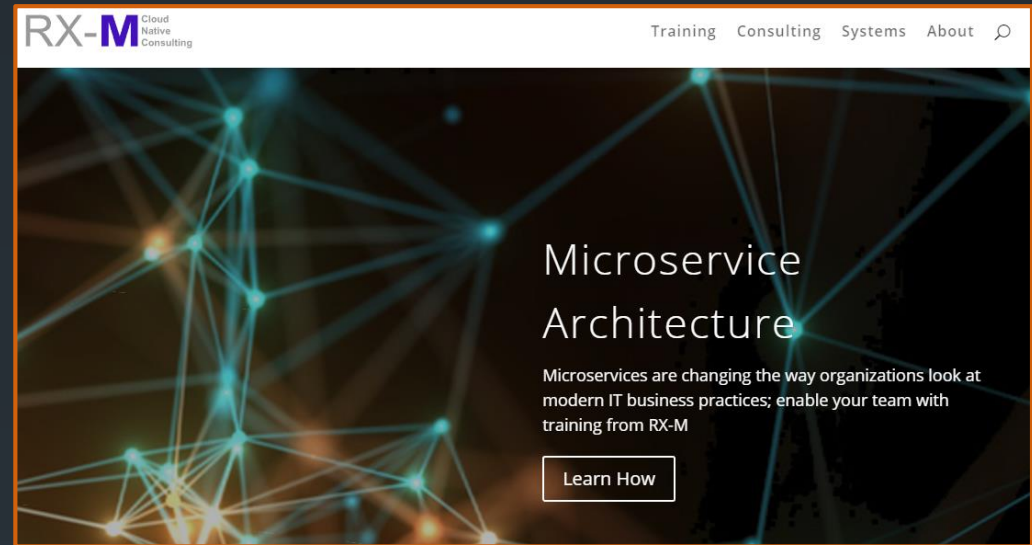
- Microservices Foundation [3 Day]
- Building Microservices on AWS [3 Day]
- Building Microservices with Go [3 Day]
- Building Microservices with Thrift [2 Day]
- Building Microservices with gRPC [2 Day]

## ▪ Container Packaged

- Docker Foundation [3 Day]
- Docker Advanced [2 Day]
- OCI [2 Day]
- CNI [2 Day]
- Containerd [2 Day]
- Rocket [2 Day]

## ▪ Dynamically Managed

- Kubernetes Foundation [2 Day]
- Kubernetes Advanced [3 Day]
- Securing Kubernetes [2 Day]
- Kubernetes Day 2 Operations [2 Day]
- Istio [2 Day]
- Knative [2 Day]



**RX-M** Cloud Native Consulting

# Overview

3

Copyright 2013-2019, RX-M LLC

## Day One

1. Containers & Orchestration
2. Kubernetes Architecture
3. Pods & Configs

## Day Two

4. Controllers
5. Services & KubeProxy
6. Managing State

## Day Three

7. Security
8. Metrics
9. Ingress

## Day Four

10. Networking
11. etcd
12. Test Prep I

## Day Five

13. Test Prep II
14. CKA Exam

## Prerequisites:

RX-M Docker Foundation  
or similar container experience

Equivalent experience:

- Familiarity with container-based microservice packaging
- Understanding of container isolation and constraints
- Basic familiarity with Docker, RunC, Rocket, Garden or similar container engine

# Administrative Info


- Course: Kubernetes CKA Boot Camp
- Length: 5 Days
- Format: Lecture/Labs/Discussion
- Schedule: 8:30AM – 5:30PM  
15 minute break, AM & PM  
1 hour lunch at noon  
Lab time at the end of each AM and PM session
- Location: Fire exits, Restrooms, Security, other matters
- Attendees: Name/Role/Experience/Goals for the Course

# Lecture and Lab

5

Copyright 2013-2019, RX-M LLC

- Our Goals in this class are two fold:
  1. **Introduce concepts and ecosystems**
    - Covering concepts and where things fit in the world is the primary purpose of the lecture/discussion sessions
    - The instructor will take you on a **tour of the museum**
      - Like a museum tour, you should interact with the instructor (tour guide), ask questions, discuss
      - Like a museum tour, you will not have time to read the slides during the tour, instead, the instructor will discuss and point out the **highlights** of the slides (exhibits) which will be waiting for you to read in depth later should you like to dig deeper
  2. **Impart practical experience**
    - This is the primary purpose of the labs
    - Classes rarely have time for complete real world projects so think of the labs as thought experiments
      - Like **hands on exhibits** at the museum



# Kubectl Drills

6

Copyright 2013-2020, RX-M LLC

## Drill 1

1. Create a pod imperatively with the `redis:latest` image and label it `app=redis` with a single replica
2. List pods with the label `app=redis`
3. Get logs for all pods with the label `app=redis`

## Drill 2

1. Set the worker node to unavailable and reschedule all the pods running on it

# Kubectl Drills

7

Copyright 2013-2020, RX-M LLC

1. Create a pod imperatively with the `redis:latest` image and label it `app=redis` with a single replica

```
kubectl run --generator=run-pod/v1 redis \
  --image=redis -l app=redis
```

2. List pods with the label `app=redis`

```
kubectl get pod -l app=redis
```

3. Get logs for all pods with the label `app=redis`

```
kubectl logs -l app=redis
```

4. Set the worker node to unavailable and reschedule all the pods running on it

```
kubectl drain nodeb (may need --ignore-daemonsets flag)
```

# Kubectl Events Drills

- Retrieve the events from the kube-system namespace
- View events generated by the Replicaset controller and Default scheduler from the whole cluster
- Retrieve events generated by all kubelets and sort them by timestamp



# Kubectl Events Drills

- Retrieve the events, if any, from the kube-system namespace

```
kubectl get events -n kube-system
```

- View events generated by the Replicaset controller and Default scheduler from the whole cluster

```
kubectl get events --all-namespaces --field-selector  
source!=kubelet -o wide
```

- Retrieve events generated by all kubelets and sort them by timestamp

```
kubectl get events --all-namespaces --field-selector  
source=kubelet --sort-by --sort-  
by=.metadata.creationTimestamp
```

# Pod Scheduling Drill

10

Copyright 2013-2020, RX-M LLC



1. Label both of your Nodes with a unique label
  - `kubectl label node master node=master`
  - `kubectl label node worker node=worker`
2. Use the simplest method to schedule a single pod on each Node

# Pod Scheduling Drill

1. Label both of your Nodes with a unique label
  - `kubectl label node ubuntu node=master`
  - `kubectl label node nodeb node=worker`
2. Use the simplest method to schedule a single pod on each Node

```
kubectl run --generator=run-pod/v1 \
  mypod --image=nginx --dry-run -o yaml
```

- Copy yaml output, pasting 2x into file (`vim mypods.yaml`)
- Place "---" between pods
- Set nodeSelector to master and worker
- Change the name of each pod
- `kubectl apply -f mypods.yaml`

Or:

```
kubectl run --generator=run-pod/v1 \
  mypod --image=nginx --overrides='{ "apiVersion": "v1",
"spec": { "nodeSelector": { "node": "master" } } }'
```

And:

```
kubectl run --generator=run-pod/v1 \
  mypod --image=nginx --overrides='{ "apiVersion": "v1",
"spec": { "nodeSelector": { "node": "worker" } } }'
```

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: mypod
    name: mypod
spec:
  nodeSelector:
    node: master
  containers:
  - image: nginx
    name: mypod
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
---
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: mypod1
    name: mypod1
spec:
  nodeSelector:
    node: worker
  containers:
  - image: nginx
    name: mypod1
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

# Deployment Drill

1. Create a Deployment with nginx version 1.14 with 3 replicas
2. Update the app to a new version 1.15.2
3. Rollback the update to the previous version

Note: make sure that you record each step

# Deployment Drill

1. Create a Deployment with nginx version 1.14 with 3 replicas

```
kubectl create deploy mydep --image=nginx:1.14  
kubectl scale deploy mydep --replicas=3
```

2. Update the app to a new version 1.15.2

```
kubectl set image deploy/mydep mydep=nginx:1.15.2 --record
```

3. Rollback the update to the previous version

```
kubectl rollout undo deploy/mydep
```

Make sure that you use the `--record` option for each step.

# Jobs Drill

1. Create a Job that:
  1. Specifies 45 completions with 2 pods running in parallel
  2. Each pod should output the string "Hello" and not get restarted if it exits successfully

# Jobs Drill

1. Create a Job that:
  1. Specifies 45 completions with 2 pods running in parallel
  2. Each pod should output the string "Hello" and not get restarted if it exits successfully

```
apiVersion: batch/v1
kind: Job
metadata:
  name: myjob
spec:
  parallelism: 2
  completions: 45
  template:
    metadata:
      name: myjob
    spec:
      containers:
      - name: myjob
        image: busybox
        command: ["sh", "-c", "echo Hello"]
        restartPolicy: OnFailure
```

# ConfigMap Drill

- Create a ConfigMap that contains this data:

```
droid=r2  
shade=light  
chastise=rebelscum
```

- Create and test a pod that mounts the ConfigMap as files with the path `/tmp/cmdata`
- The Pod should cat one of the files and exit



# ConfigMap Drill

17

Copyright 2013-2020, RX-M LLC

- Create a ConfigMap that contains this data:

```
droid=r2
shade=light
chastise=rebelscum
```

```
kubectl create cm mycm \
--from-literal=droid=r2 \
--from-literal=shade=light \
--from-literal=chastise=rebelscum
```

- Create and test a pod that mounts the ConfigMap as files with the path /tmp/cmdata
- The Pod should cat one of the files and exit when it runs
  - `kubectl run \`  
  `--generator=run-pod/v1 \`  
  `cmpod --image=busybox \`  
  `--dry-run \`  
  `-o yaml > cmpod.yaml`
  - Edit the file, adding the CM

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: cmpod
  name: cmpod
spec:
  containers:
  - image: busybox
    name: cmpod
    command: ["sh", "-c", "cat /tmp/cmdata/chastise"]
    volumeMounts:
    - mountPath: "/tmp/cmdata"
      name: data
  dnsPolicy: ClusterFirst
  restartPolicy: Never
  volumes:
  - name: data
    configMap:
      name: mycm
```

# Secrets Drill

1. Create a secret that sets a username and password *from literals*
2. Create a Pod that uses the secret as environment variables
3. Check your work by exec-ing into the pod, getting the values of the env vars

# Secrets Drill

1. Create a secret that sets a username and password *from literals*

```
kubectl create secret generic mysec \  
--from-literal=username=user \  
--from-literal=password=pass
```

2. Create a Pod that uses the secret as environment variables

```
kubectl run --generator=run-pod/v1 \  
--env=USERNAME=user,PASSWORD=pass \  
secpod --image=nginx --dry-run \  
-o yaml > secpod.yaml
```

Edit the yaml, or omit the --env= flag and use envFrom in the yaml

3. Check your work by exec-ing into the pod, getting the values of the env vars

```
kubectl exec -it secpod -- /bin/bash  
$ echo $USERNAME $PASSWORD
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  creationTimestamp: null  
  labels:  
    run: secpod  
  name: secpod  
spec:  
  containers:  
  - image: nginx  
    name: secpod  
    env:  
    - name: USERNAME  
      valueFrom:  
        secretKeyRef:  
          name: mysec  
          key: username  
    - name: PASSWORD  
      valueFrom:  
        secretKeyRef:  
          name: mysec  
          key: password  
--- OR ---  
  envFrom:  
  - secretRef:  
    name: mysec  
--- \OR ---  
  resources: {}  
  dnsPolicy: ClusterFirst  
  restartPolicy: Always  
status: {}
```

# Namespaces Drill

1. Create a namespace named `myns`
2. Apply a quota:
  - `PersistentVolumeClaims`: 1
  - `nodePorts`: 2
  - `Pods`: 5
3. Create a Role in your new namespace that allows the verbs "get" and "list" on Pods
4. Bind that role to a user named Paul

# Namespaces Drill

21

Copyright 2013-2020, RX-M LLC

## 1. Create a namespace named myns

```
kubectl create ns myns
```

## 2. Apply a quota:

- Persistentvolumeclaims: 1
- nodePorts: 2
- Pods: 5

```
kubectl create quota myquota \  
--hard=count/pvc=1,count/service.nodeport=2,count/pods=5 \  
--namespace=myns
```

## 3. Create a Role in your new namespace that allows the verbs "get" and "list" on Pods

```
kubectl create role myrole --verb=get,list --resource=pod -n myns
```

## 4. Bind that role to a user named Paul

```
kubectl create rolebinding myrb --role=myrole --user=paul -n myns
```

# Network Policy Drill

- Run a server pod in the default namespace

```
kubectl run server --generator=run-pod/v1 --image rxmllc/hostinfo --port 80 -l app=info
```

- Run a client pod in another namespace called drill

```
kubectl create ns drill -l purpose=test
```

```
kubectl label ns drill purpose=test
```

```
kubectl run client -n drill --generator=run-pod/v1 --image busybox --command -- tail -f /dev/null
```

- Adjust your cluster so the client pod in the drill namespace can communicate with the server pod while a default-deny network policy is in place. Do not delete the network policy:

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny
spec:
  podSelector: {}
  policyTypes:
  - Ingress
```

# Network Policy Drill

1. Adjust your cluster so the client pod in the drill namespace can communicate with the server pod while a default-deny network policy is in place. Do not delete the network policy:

Create the following permissive network policy:

nano np-client.yaml:

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: access-hostinfo
spec:
  podSelector:
    matchLabels:
      app: info
  ingress:
  - from:
    - namespaceSelector:
        matchLabels:
          purpose: "test"
```

```
kubectl apply -f np-client.yaml
```

```
kubectl exec -it client -- wget -qO - http://<server pod IP>:9898
```

# Kubelet Drill

1. Create a PodSpec that runs a single container
  - Deploy the pod w/o using the api-server



# Kubelet Drill

1. Create a PodSpec that runs a single container
  - Deploy the pod w/o using the api-server

```
sudo cat /var/lib/kubelet/config.yaml | grep -i manifest
```

```
kubect1 run --generator=run-pod/v1 kubeletpod --image=nginx \  
--dry-run=true -o yaml
```

- copy the file to /etc/kubernetes/manifests (or the manifest path)

# Persistent Storage Drill

1. Statically provision a PV, create a matching PVC and a Pod that uses the PVC as a volume mount

# Persistent Storage Drill

1. Statically provision a PV, create a matching PVC and a Pod that uses the PVC as a volume mount

```
$ mkdir /home/user/pv
```

```
$ kubectl run --generator=run-pod/v1 \
kubetepod --image=nginx \
--dry-run=true -o yaml
```

```
$ kubectl exec -it persistentpod bash
```

```
root@persistentpod:/# touch /var/www/html/file
```

```
root@persistentpod:/# exit
```

```
$ ls -l pv
```

```
-rw-r--r-- 1 root root 0 Aug 17 12:16 file
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: hppv
spec:
  capacity:
    storage: 100Mi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Retain
  hostPath:
    path: /home/user/pv
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: hppvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 100Mi
```

```
apiVersion: v1
kind: Pod
metadata:
  name: persistentpod
spec:
  containers:
    - name: persistentcontainer
      image: nginx
      volumeMounts:
        - mountPath: "/var/www/html"
          name: persistence
  volumes:
    - name: persistence
      persistentVolumeClaim:
        claimName: hppvc
```

# Shared Storage Drill

1. Create a pod with the following attributes:
  - a. The pod must run two busybox containers
    - i. Configure both busybox containers to run the `tail -f /dev/null` command
  - b. The pod must have a shared host volume which will be removed by Kubernetes when the pod is deleted
  - c. Create a file in the volume called “message”; make sure that this file is created before either of the two main containers run
  - d. Mount the volume at `/vola` in one container and at `/data` in the other
  - e. Exec into both containers to verify the “message” file is accessible on the appropriate path in both containers

# Shared Storage Drill

1. Create a pod with the following attributes:
  - a. The pod must run two busybox containers
    - i. Configure both busybox containers to run the “tail -f /dev/null” command
  - b. The pod must have a shared host volume which will be removed by Kubernetes when the pod is deleted
  - c. Create a file in the volume called “message”; make sure that this file is created before either of the two main containers run
  - d. Mount the volume at “/vola” in one container and at “/data” in the other
  - e. Exec into both containers to verify the “message” file is accessible on the appropriate path in both containers

```
kubect1 run --generator=run-pod/v1 \  
kubect1 run --image=nginx \  
--dry-run=true -o yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: drillpod  
spec:  
  restartPolicy: Always  
  volumes:  
    - name: shared-vol  
      emptyDir: {}  
  initContainers:  
    - name: msg-creator  
      image: busybox  
      command: ["sh", "-c"]  
      args: ["touch /share/message"]  
  volumeMounts:  
    - mountPath: /share  
      name: shared-vol  
  containers:  
    - name: msg-one  
      image: busybox  
      command: ["tail", "-f", "/dev/null"]  
      volumeMounts:  
        - mountPath: /vola  
          name: shared-vol  
    - name: msg-two  
      image: busybox  
      command: ["tail", "-f", "/dev/null"]  
      volumeMounts:  
        - mountPath: /data  
          name: shared-vol
```

# Etcd Drill

1. Backup etcd by copying the member/snap/db file
2. Backup an etcd cluster using etcdctl
3. Save the backups to /tmp/cka

# Etcd Drill

31

Copyright 2013-2020, RX-M LLC

1. Backup etcd by copying the member/snap/db file

```
sudo cp /var/lib/etcd/member/snap/db /tmp/cka/
```

2. Backup an etcd cluster using etcdctl

```
kubectl -n kube-system exec etcd-labsys -- \
/bin/sh -c "ETCDCTL_API=3 etcdctl \
--endpoints=https://127.0.0.1:2379 \
--cacert=/etc/kubernetes/pki/etcd/ca.crt \
--cert=/etc/kubernetes/pki/etcd/server.crt \
--key=/etc/kubernetes/pki/etcd/server.key \
snapshot save /var/lib/etcd/backup.db"
```

3. Save the backups to /tmp/cka

```
sudo cp /var/lib/etcd/backup.db /tmp/cka
```

# Add Node Drill

## 1. Add a worker node with kubeadm

- Install docker:
  - `wget -O - https://get.docker.com | sh`
- Add the Kubernetes Apt Repo:
  - `curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -`
  - `echo "deb http://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee -a /etc/apt/sources.list.d/kubernetes.list`
- Use the quickest way to join a node to your cluster



# Add Node Drill

## 1. Add a worker node with kubeadm

- Install docker:
  - `wget -O - https://get.docker.com | sh`
- Add the Kubernetes Apt Repo:
  - `curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -`
  - `echo "deb http://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee -a /etc/apt/sources.list.d/kubernetes.list`
- Use the quickest way to join a node to your cluster
  - On Master
    - `kubeadm token create --print-join-command`
  - On worker
    - `kubeadm join master:6443 \`  
`--token 1ab2c3.v7zjj0oj39ypq0x8 \`  
`--discovery-token-ca-cert-hash sha256:...`

# Kubelet Service Drill

1. Identify the following settings for one of your Kubernetes nodes
2. Make the cka\_kubelet service persists through reboots
3. Enable server certificate rotation on the Kubelet

# Kubelet Service Drill

## 1. Identify the following settings for one of your Kubernetes nodes

- Kubelet Service Status
- Path to Kubelet Systemd Service
- API Server IP Address
- Kubelet Static Pod Manifest Path

```
sudo systemctl status kubelet
sudo systemctl status kubelet | grep "Loaded:"
sudo cat /etc/kubernetes/kubelet.conf | grep "server:"
sudo cat /var/kubelet/config.yaml | grep "static"
```

## 2. Make the cka\_kubelet service persists through reboots

```
sudo systemctl enable cka_kubelet
```

## 3. Enable server certificate rotation on the Kubelet

Add the flag --rotate-server-certificates to the end of that file:

```
sudo nano /etc/systemd/system/kubelet.service.d/10-kubeadm.conf
```

Or: append --rotate-server-certificates to the end of the kubelet systemd service ExecStart

```
sudo nano /lib/systemd/system/kubelet.service
```

Then reload the systemd service:

```
sudo systemctl daemon-reload
sudo systemctl restart kubelet
```

# Kubelet Config Drill

1. Create a kubelet.conf file that connects to your cluster as a node named "new\_worker"
2. Set up the new kubelet.conf to use certificates found /tmp/

# Kubelet Config Drill

1. Create a kubelet.conf file that connects to your cluster as a node named "new worker"

```
sudo cat /etc/kubernetes/kubelet.conf  
sudo cp /etc/kubernetes/kubelet.conf /tmp/cka_kubelet.conf  
nano /tmp/cka_kubelet.conf
```

Change all references to `system:node:ubuntu` to `system:node:new_worker`

```
apiVersion: v1  
clusters:  
- cluster:  
  certificate-authority-data: ...  
  server: https://<Cluster API Server>:6443  
  name: kubernetes  
contexts:  
- context:  
  cluster: kubernetes  
  user: system:node:new_worker  
  name: system:node:new_worker@kubernetes  
current-context: system:node:new_worker@kubernetes  
kind: Config  
preferences: {}  
users:  
- name: system:node:new_worker  
  user:  
    client-certificate: /var/lib/kubelet/pki/kubelet-client-current.pem  
    client-key: /var/lib/kubelet/pki/kubelet-client-current.pem
```

# TLS Bootstrap Drill

38

Copyright 2013-2020, RX-M LLC

- Use Kubectl to generate a bootstrap-kubeconfig file under /tmp/bootstrap-kubeconfig.conf
  - Find out your current connection information
  - Your CA certificate is found under /etc/kubernetes/pki/ca.crt
  - Try to find a bootstrap token in your cluster
    - If you do not have one, use kubeadm to create a token

# TLS Bootstrap Drill

Use Kubectl to create a bootstrap-kubeconfig file under /tmp/bootstrap-kubeconfig.conf

- Find out your current connection information

```
kubectl config view
```

- Your CA certificate is found under /etc/kubernetes/pki/ca.crt

```
kubectl config set-cluster bootstrap \  
--kubeconfig=/tmp/bootstrap-kubeconfig \  
--server='https://172.31.6.52:6443' \  
--certificate-authority=/etc/kubernetes/pki/ca.crt
```

- Try to find a bootstrap token in your cluster

```
kubectl get secrets -n kube-system | grep bootstrap
```

- Use base64 to decode the token-id and token-secret

```
kubectl get secret -n kube-system $BOOTSTRAPSECRET -o jsonpath='{.data.token-id}' | base64 --  
decode  
kubectl get secret -n kube-system $BOOTSTRAPSECRET -o jsonpath='{.data.token-secret}' | base64  
--decode
```

- If you do not have one, use kubeadm to create a token

```
kubeadm token create
```

# TLS Bootstrap Drill (cont.)

- Associate the token with a user:

```
kubectl config set-credentials kubelet-bootstrap \  
--kubeconfig=/tmp/bootstrap-kubeconfig \  
--token=ppr919.z8k9u58ng23hqe5m
```

- Associate a user with a cluster using a context:

```
kubectl config set-context bootstrap \  
--kubeconfig=/tmp/bootstrap-kubeconfig \  
--user=kubelet-bootstrap \  
--cluster=bootstrap
```

- Create the file by switching the context:

```
kubectl config use-context bootstrap \  
--kubeconfig=/tmp/bootstrap-kubeconfig
```