

# Kubernetes

## Lab 2 – Kubernetes Exploration

Kubernetes clusters track and manage objects of various “kinds”. Applications make use of four kinds of objects in particular:

- **Pods** – groups of containers deployed as a unit
- **Replica Sets** – sets of pods defined by a template which the Controller Manager replicates across the cluster
- **Deployments** – a rollout strategy for pods and replica sets
- **Services** – end points used to distribute requests to one of a pod’s replicas

Thus basic Kubernetes applications consist of pods, which implement the application functionality; replica sets, which ensure pods are always available; and Services which expose a dynamic set of pods to clients as a single endpoint. Deployments describe how to launch or upgrade a given application (set of pods).

The principal tool used to manage kubernetes clusters is kubectl. Using kubectl you can create deployments and services, monitor cluster components and pods, upgrade deployments and remove resources that are no longer required.

In this lab we will get familiar with kubectl and learn how to use it to manage many aspects of Kubernetes.

### 1. kubectl

The `kubectl` command provides a range of features we can use with Kubernetes. Run `kubectl` without arguments to get a list of the available commands.

```
user@ubuntu:~$ kubectl

kubectl controls the Kubernetes cluster manager.

Find more information at:
https://kubernetes.io/docs/reference/kubectl/overview/

Basic Commands (Beginner):
  create          Create a resource from a file or from stdin.
  expose          Take a replication controller, service, deployment or pod and
expose it as a new Kubernetes Service
  run             Run a particular image on the cluster
  set             Set specific features on objects

Basic Commands (Intermediate):
  explain         Documentation of resources
  get            Display one or many resources
  edit           Edit a resource on the server
  delete         Delete resources by filenames, stdin, resources and names, or
by resources and label selector

Deploy Commands:
  rollout        Manage the rollout of a resource
  scale          Set a new size for a Deployment, ReplicaSet, Replication
Controller, or Job
  autoscale      Auto-scale a Deployment, ReplicaSet, or ReplicationController

Cluster Management Commands:
  certificate     Modify certificate resources.
  cluster-info    Display cluster info
  top            Display Resource (CPU/Memory/Storage) usage.
  cordon         Mark node as unschedulable
  uncordon       Mark node as schedulable
  drain          Drain node in preparation for maintenance
  taint          Update the taints on one or more nodes

Troubleshooting and Debugging Commands:
```

```

describe      Show details of a specific resource or group of resources
logs          Print the logs for a container in a pod
attach        Attach to a running container
exec          Execute a command in a container
port-forward  Forward one or more local ports to a pod
proxy         Run a proxy to the Kubernetes API server
cp            Copy files and directories to and from containers.
auth          Inspect authorization

Advanced Commands:
diff          Diff live version against would-be applied version
apply         Apply a configuration to a resource by filename or stdin
patch         Update field(s) of a resource using strategic merge patch
replace       Replace a resource by filename or stdin
wait          Experimental: Wait for a specific condition on one or many
resources.
convert       Convert config files between different API versions
kustomize     Build a kustomization target from a directory or a remote url.

Settings Commands:
label         Update the labels on a resource
annotate      Update the annotations on a resource
completion    Output shell completion code for the specified shell (bash or
zsh)

Other Commands:
api-resources  Print the supported API resources on the server
api-versions  Print the supported API versions on the server, in the form of
"group/version"
config        Modify kubeconfig files
plugin        Provides utilities for interacting with plugins.
version       Print the client and server version information

Usage:
  kubectl [flags] [options]

Use "kubectl <command> --help" for more information about a given command.
Use "kubectl options" for a list of global command-line options (applies to all
commands).

user@ubuntu:~$

```

Take a moment to review available options. One useful subcommand is the global options, take a moment to review the output of `kubectl options`.

To use the `kubectl` command to control a remote cluster we must specify the cluster endpoint to `kubectl`. The `kubectl` command can be used to control several clusters from a single workstation. Clusters are given a name and settings, including the IP address and port of the cluster API service.

To get configuration help issue the `kubectl help` subcommand.

```

user@ubuntu:~$ kubectl help config

Modify kubeconfig files using subcommands like "kubectl config set current-context my-context"

The loading order follows these rules:

1. If the --kubeconfig flag is set, then only that file is loaded. The flag may only be set
once and no merging takes place.
2. If $KUBECONFIG environment variable is set, then it is used as a list of paths (normal
path delimiting rules for your system). These paths are merged. When a value is modified, it is modified in the file that
defines the stanza. When
a value is created, it is created in the first file that exists. If no files in the chain exist,
then it creates the
last file in the list.
3. Otherwise, ${HOME}/.kube/config is used and no merging takes place.

```

#### Available Commands:

current-context	Displays the current-context
delete-cluster	Delete the specified cluster from the kubeconfig
delete-context	Delete the specified context from the kubeconfig
get-clusters	Display clusters defined in the kubeconfig
get-contexts	Describe one or many contexts
rename-context	Renames a context from the kubeconfig file.
set	Sets an individual value in a kubeconfig file
set-cluster	Sets a cluster entry in kubeconfig
set-context	Sets a context entry in kubeconfig
set-credentials	Sets a user entry in kubeconfig
unset	Unsets an individual value in a kubeconfig file
use-context	Sets the current-context in a kubeconfig file
view	Display merged kubeconfig settings or a specified kubeconfig file

#### Usage:

kubectl config SUBCOMMAND [options]

Use "kubectl <command> --help" for more information about a given command.

Use "kubectl options" for a list of global command-line options (applies to all commands).

user@ubuntu:~\$

Run the `kubectl config view` subcommand again to display the current client configuration.

```
user@ubuntu:~$ kubectl config view
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: DATA+OMITTED
    server: https://192.168.228.157:6443
    name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: kubernetes-admin
    name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    client-certificate-data: REDACTED
    client-key-data: REDACTED

user@ubuntu:~$
```

When you run *kubectl* commands a context is required. The context tells *kubectl* which cluster to connect to and which user to authenticate as. As you can see the values kubeadm configured means the `kubectl` command tries to reach the API server on port 6443 via our host's IP with TLS.

To view the REDACTED elements, add `--flatten`.

We can configure `kubectl` explicitly so that we can adjust our cluster settings in the future if need be. Get help on the `config set-cluster` subcommand:

```
user@ubuntu:~$ kubectl help config set-cluster
```

Sets a cluster entry in kubeconfig.

Specifying a name that already exists will merge new fields on top of existing values for those fields.

#### Examples:

```
# Set only the server field on the e2e cluster entry without touching other values.
kubectl config set-cluster e2e --server=https://1.2.3.4
```

```
# Disable cert checking for the dev cluster entry
kubectl config set-cluster e2e --insecure-skip-tls-verify=true
```

```
--embed-certs=false: embed-certs for the cluster entry in kubeconfig
```

```
kubectl config set-cluster NAME [--server=server] [--certificate-
authority=path/to/certificate/authority]
[--insecure-skip-tls-verify=true] [options]
```

```
user@ubuntu:~$
```

```
user@ubuntu:~$ ls -la ~/.kube/
```

```
user@ubuntu:~$
```

```
user@ubuntu:~$ cat ~/.kube/config
```

Page 4/18 © Copyright Monday, Jan 27, 2020, 7:50 AM RX-M LLC

```
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSB0tLS0tCk1JSUM4akNDQWRxZ0F3SUJBZ0lJWjNqenAweEZVaDR3RFFZSk1vWk1odmNOQV
FFTEJRQXQdGVEVUTUJFR0ExVUUQKXhNS2EzVmlaWEp1WlhSbGN6QWVGdzB4T1RBNE1EWXlNREl6TURkYUz3MHlNREE0TURVeU1E
SXpNRGxhTURReApGekFWQmdOVk1Bb1REbK41YzNSbGJUCkRZWE4wWlhKek1Sa3dGd1lEVlFRREV4QnJkV0psY201bGRHVnpMV0
ZrCmJXbHVNSU1CSWpBTk1Jna3Foa2lHOXcwQkFRRUZBQ0U9DQVE4QU1JSUJDZ0tDQVFFQXJ3anFoQVBJdkFVTWh0Nk8kdWFMWEZz
b0ZPTWJhbGh1JTkZmZy9nb3FwTE93UWFGRUVjUV1UeXdhSUZvbERHekZ5eko4VFJQYjK5cld0K1IvTgoxeEVHODkwdnltc1Q5Tm
80U042REI4OTZtRWQ3djdBbTgxaWY5NWlQVG5SbkpFcxWwMuk50XdsRyTaR205VW1QC1p3ZVRTL05tb3FTRFdvYjAzUnFrVXg2
bDdoWW91c1pzSkRiCvJHUWwrcWxtb0x1K1NVUfDGB0FwYjF0SDG0dS8KZjRXL25RQzFqS0U3bEJ6UzZRSXhIMkNHbUo1MXRydv
dlldjZVYXW6TW1iUy8wVi9KT0xLSW90SVQRZG9oZHpqZQ05a0t5dDNJ03Yxakg2Ukh5cwo0RED6TmFQby91Z3VTL2RNT3BnNi9r
ZHD0XpXegJSb2RPUUXFZEVGZFRWRkxiCmY5ZWVLD0lEQVFBQm95Y3dKVEFPQmdOVkhROEJBZjhFQkFNQ0JhQXdFdl1EVlIiwB
EBD3dDZ1lJS3dZQk1RVUgKQXZjD0RRWUplb1pJaHJzTkFRRUxUUFEZ2dFQkFDVGxLVWI4SXdXSFl4T0RrVgTvwWlun1ZtNEdB
T3pHdk84RwpaTC9UUVjU2YnN0cGtMUWRWmg44RGVrODJRrRHBWnm9wN2dFejc2TVJHMHj4Q1lTODkxeFgyVS9PY0kwVThDVFpCm
pSams00V1YTWVMUzVSTHp4dW1Ub3RHMlVQbVhydHRvanhBR3ZoaH55R3RTR1ZiClBKbk3ZEZ2T1h4elhh0TUKNWZxV3YyV21T
UXZ0cUFJL2d0L1c0bmg4QU5iKzEvYjIrb3lQ01pGcFVqSGhPb21yV05LNjErL1JXVjU5K3pDaQpRZ1VaeGNLU1Q0SszJFNGxkaG
5SdzJNAFFsv3dJckR4dGNKbGtseDZMTUUhZ0w2WnFLUKVsU1IvaWJNB25WU0pzCk1JYk5tNjB4cTlpaWs5SUVBWNyXUERXcFB5
b0M5MVRTSsys3OUY2TEF4T1JNejdLSjZTWT0KLS0tLS1FTkQgQ0VSVElGSUNBVEU0tLS0tLQo=
```

client-key-data:

```
LS0tLS1CRUdJTiB0EUFJJKFURSB0tLS0tLQpNSU1Fb2dJQkFBS0NB0UUVBcndqcWhBUEl2QVNaG0T23VhTFhGc29GT0
1iYWxoY05GZmcvZ29xcExPd1FhRkVFCmNRWVR5d2FJRm9sRed6Rn16SjhUUBiOT1yV3QrUi90MXhFRzg5MHZ5bXNUOU5vNFNO
NkRCODk2bUVkN3Y3QW0KODFPzjk1aVBUB1JuSkVxbHAXSTk5d1JHK1pHbTlVbVBad2VUUY90bW9xU0RXb2IwM1Jxa1V4Nm3aF
lvdXNacwPKREhXUkdRaGtXbG1VTGUU1VQV0ZvQXBiMXRIODR1L2Y0cS9uUUMxaktFN2xCe1M2U1L4SDJDR21KNTF0cnVXCmV2
N1VhbHpnBwJTLzBWL0pTETtJb3RJVCtkb2hkem1OWtLeXQzSVN2MwPwIN1JTeXFqNERHek5hUG8vdWd1Uy8KZE1PcGc2L2tkd0
Q5eld4Y1JvZE9RTEVkrUZkVFZGTGJmOWVlS3dJREFRQUJbB0lCQUFuNFV3dGovOG9knjRHNAp5RzV3aFJucU5mUHU50EoyMjZy
MXNnQm9qNhhjZ2U2L05xYU5keEVvZFJYN2laUUI5QTB0dnBZK0QyYy9JQUxCCnVnTFhHSE5KbFJCTm14eG1WYk1vNUVQTW9Lb0
40RkdoNgdZWWRUd3d0M0Y3bGVodVRCTkV6dnPQkpyd25mYjQKME5R3NqQXdwQmROa2FIVmNCc3NDbElsdThQYzdTSTBuK0ty
ZzJ6akpYcU4rMHQxYVdYMWdXb2RMclJ1UWkzTQpFbXZ2SUNpazBHNnB0Zmh6K2R6dC9ReWJmVz1DWXRrdk9L3dPWkpvm1NscT
lzMGLQYkZBMGo3d1B1Ti8zNW80CmZTUkRwQkh2VmxGTEl0bG9gMXkxTWYwTEl0U52NGJzNWFW0Wmh00FZxdElSanYyZm5BSWxK
MnhPb2E5dHJGMDgKeE0cnBhRUNnWUVBMnJqNWNCaVovQjRaYj1zeERXblNFRGt0RHNYTA2cEpnRzRQOXFXTlYvbGk1Vmg5MW
NkApQa1l6ZEFncXRHUGRMbG5vWEVYaE9sd3pwOU8ydzh0Ukhsd1tnZ0s5K1QzZzNWwjdXYkpxZTlWbTBFaDNVaFRwCm5pZnhX
bn1HZkRaR1JpMG52WHRjVEp6MVU5Q0Z3bVA4MVlBUFU4T1BLTkxZbCtuSUd6cmw2anNDZ1lFQXpOM1EKYUQyUEPwTzRMLzBMZF
NuMHh0ek1zRytcUJkYXRMZDFtUXpYNTFpcHZAwwdNeHdNcENRV0lSKz1Ra2RybmPzUwozS3h1ZGc5ZGYybmlRNSsttbENFS01a
S1VUaDNJRnNxc1VwcG50TD1pwjRXR2xkR1JXdS9aSUZjQTIvSHZvSGlMCjd4azExYkZvTmJ3T3EzU0JDUHnNqMvQZFBXaUpnMU
5zSXFwUDd0RURnnWUJXcWxTC0JoTj13cT14a3Zpc3gvRmsKWEVODVJ4ZWVKeXRZcUUVQTEENXbFlJb2o5TjBNbElnNmhqazkwK014
RFRnNmweG9DOWUxNG9uQVZYOTViVlZSUApJMFNGWNESEfHm0lCaXg2TG1malNYdWpyQk1iZ3czT2pTVWxKwKprU1YwekVWRm
Qwek9QWitJTMcrSjhWQUlXcJFzTTcvWkVNSy8rR3NpdU1IcDVlWVFLQmdHSisvdFV2UzZSaHBQdUZpTDJ6QjFtWkROeURSQitU
NHV0aURTc2gKanFoTzY0VV1LYkVJK2xic2RxdEVURVVTZTM1Y2R5TWIwQnY0OVRyDUhYznZ5VElNMUK3UzBLK21CL0pWVEp0eA
pXdlhxNGcVdGxiQndLOU14NE0xNHB4UUVlOT0tOTXBjCEo4WUUVckJmRXhxQjhBdjJXUV1lV0VoTysxWmxoR2NDC1JWMFJBb0dB
YXlsTmPCL1VENzRoUzFVbCtZ25wMulaRgd1UkFpR0JyUvAYZ0p1SUthREFUN3dkOG5uamdzMXIKc3hHWlFQZlFOY2tRTXFXVeT
R5M0c2Y1ZaNUp00UQRmNPS01IcDZFa1I2SENKbVBFV1dMLzNEek1hbKZrSHE10Ao0L08zV3ptT2d1VXI5MlJyc0MrMGfPMGw0
N1hSNnB1ZndE0EcZNYtMY3U5UnJxelpSSTA9Ci0tLS0tRU5EIFJTQSBQUklWQVRFIEtFWS0tLS0tCg==
```

```
user@ubuntu:~$
```

The `kubectl config view` command will display nearly the same data, obfuscating the key data. The config file is simple and can easily be pre-generated and distributed to any client systems that require connection to a given cluster.

## 2. Test the Cluster

Now with our cluster running and `kubectl` configured lets issue some commands to test the Kubernetes cluster. The `cluster-info` subcommand can be used to test the cluster API end point and the `get nodes` command can be used to see the nodes in the cluster.

```
user@ubuntu:~$ kubectl cluster-info
```

```
Kubernetes master is running at https://192.168.228.157:6443
KubeDNS is running at https://192.168.228.157:6443/api/v1/namespaces/kube-system/services/kube-
dns:dns/proxy
```

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

```
user@ubuntu:~$
```

If you are really adventurous run the suggested command for a detailed cluster overview, careful though, its a lot of information!

```
user@ubuntu:~$ kubectl cluster-info dump |& wc -l
```

```
4972
```

```
user@ubuntu:~$
```

To get detailed node information use the `describe node` subcommand again on the desired node name:

```
user@ubuntu:~$ kubectl describe node ubuntu
```

```
Name:                ubuntu
Roles:               master
Labels:              beta.kubernetes.io/arch=amd64
                    beta.kubernetes.io/os=linux
                    kubernetes.io/arch=amd64
                    kubernetes.io/hostname=ubuntu
                    kubernetes.io/os=linux
                    node-role.kubernetes.io/master=
Annotations:         kubeadm.alpha.kubernetes.io/cri-socket: /var/run/dockershim.sock
                    node.alpha.kubernetes.io/ttl: 0
                    volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp:   Wed, 08 Jan 2020 12:15:29 -0800
Taints:              <none>
Unschedulable:       false
Conditions:
  Type                 Status  LastHeartbeatTime               LastTransitionTime
Reason                Message
-----
NetworkUnavailable    False   Wed, 08 Jan 2020 12:31:56 -0800   Wed, 08 Jan 2020 12:31:56 -0800
WeaveIsUp              Weave pod has set this
MemoryPressure         False   Wed, 08 Jan 2020 12:40:33 -0800   Wed, 08 Jan 2020 12:15:27 -0800
KubeletHasSufficientMemory kubelet has sufficient memory available
DiskPressure           False   Wed, 08 Jan 2020 12:40:33 -0800   Wed, 08 Jan 2020 12:15:27 -0800
KubeletHasNoDiskPressure kubelet has no disk pressure
PIDPressure            False   Wed, 08 Jan 2020 12:40:33 -0800   Wed, 08 Jan 2020 12:15:27 -0800
KubeletHasSufficientPID kubelet has sufficient PID available
Ready                  True    Wed, 08 Jan 2020 12:40:33 -0800   Wed, 08 Jan 2020 12:32:02 -0800
KubeletReady           kubelet is posting ready status. AppArmor enabled
Addresses:
  InternalIP: 192.168.228.157
  Hostname:   ubuntu
Capacity:
  cpu:                2
  ephemeral-storage:  18447100Ki
  hugepages-1Gi:      0
  hugepages-2Mi:      0
  memory:              2030628Ki
  pods:               110
Allocatable:
  cpu:                2
  ephemeral-storage:  17000847332
  hugepages-1Gi:      0
  hugepages-2Mi:      0
  memory:              1928228Ki
  pods:               110
System Info:
  Machine ID:          6e883acc04fc7db3713776be57a3dac9
  System UUID:         5FBB4D56-33A0-3A9A-19B1-95D19AECC42F
  Boot ID:             7c1dbf59-0da5-4010-8ea0-b7253b5446e4
  Kernel Version:      4.4.0-31-generic
  OS Image:            Ubuntu 16.04.1 LTS
  Operating System:    linux
  Architecture:        amd64
  Container Runtime Version: docker://19.3.5
  Kubelet Version:     v1.16.4
  Kube-Proxy Version:  v1.16.4
Non-terminated Pods:  (8 in total)
  Namespace           Name
  ----
  kube-system         coredns-5644d7b6d9-b4rnz
(3%)                  170Mi (9%)      25m
CPU Requests  CPU Limits  Memory
```

```

kube-system          coredns-5644d7b6d9-lxdqv      100m (5%)    0 (0%)    70Mi
(3%)          170Mi (9%)    25m
kube-system          etcd-ubuntu                    0 (0%)      0 (0%)    0 (0%)
0 (0%)          24m
kube-system          kube-apiserver-ubuntu          250m (12%)   0 (0%)    0 (0%)
0 (0%)          24m
kube-system          kube-controller-manager-ubuntu 200m (10%)   0 (0%)    0 (0%)
0 (0%)          24m
kube-system          kube-proxy-npxks              0 (0%)      0 (0%)    0 (0%)
0 (0%)          25m
kube-system          kube-scheduler-ubuntu         100m (5%)    0 (0%)    0 (0%)
0 (0%)          24m
kube-system          weave-net-rvhvk               20m (1%)     0 (0%)    0 (0%)
0 (0%)          9m42s
Allocated resources:
(Total limits may be over 100 percent, i.e., overcommitted.)
Resource           Requests    Limits
-----
cpu                 770m (38%) 0 (0%)
memory              140Mi (7%) 340Mi (18%)
ephemeral-storage   0 (0%)     0 (0%)
Events:
Type      Reason              Age           From          Message
----      -
Normal    NodeHasSufficientMemory 26m (x8 over 26m) kubelet, ubuntu Node ubuntu status is
now: NodeHasSufficientMemory
Normal    NodeHasNoDiskPressure 26m (x8 over 26m) kubelet, ubuntu Node ubuntu status is
now: NodeHasNoDiskPressure
Normal    NodeHasSufficientPID    26m (x7 over 26m) kubelet, ubuntu Node ubuntu status is
now: NodeHasSufficientPID
Normal    Starting              25m           kube-proxy, ubuntu Starting kube-proxy.
Normal    NodeReady              9m27s         kubelet, ubuntu Node ubuntu status is
now: NodeReady

user@ubuntu:~$

```

Describe provides a wealth of node information. Your report will be similar but different than the one above.

- How much memory does your node have?
- How many CPUs?
- How many pods can your node run?
- What container runtime is the `kubelet` using?
- What version of `kubelet` is your node running?

Previously we used the `version` subcommand to discover the version of the `kubect1` client but now that our config is in place we can also see the version of the cluster API Server.

```

user@ubuntu:~$ kubectl version

Client Version: version.Info{Major:"1", Minor:"16", GitVersion:"v1.16.4",
GitCommit:"224be7bdce5a9dd0c2fd0d46b83865648e2fe0ba", GitTreeState:"clean", BuildDate:"2019-12-
11T12:47:40Z", GoVersion:"go1.12.12", Compiler:"gc", Platform:"linux/amd64"}
Server Version: version.Info{Major:"1", Minor:"16", GitVersion:"v1.16.4",
GitCommit:"224be7bdce5a9dd0c2fd0d46b83865648e2fe0ba", GitTreeState:"clean", BuildDate:"2019-12-
11T12:37:43Z", GoVersion:"go1.12.12", Compiler:"gc", Platform:"linux/amd64"}

user@ubuntu:~$

```

If you are familiar with Golang, notice the use of the `gc` tool chain (vs `gccgo`).

### 3. Creating Applications

With our cluster running and `kubect1` configured we can try to start a simple application on the cluster. The `kubect1` command provides a `get` subcommand which can be used to get information on any one of the key Kubernetes component types: deployments, pods, replica sets, and Services. While you can type `kubect1 get replicaset`, that would be fairly inhumane so `kubect1` allows you to use the abbreviation `rs` for replica sets.

If you want to save yourself even more typing you can take advantage of kubectl's tab completion functionality. Try typing kubectl and then :

That is, press the tab key twice...

```
user@ubuntu:~$ kubectl get
Desktop/                  .kube/                  Public/
...
user@ubuntu:~$ kubectl get
```

This output is the standard bash shell completion, which just lists the files in the working directory. Not very helpful. You can enable temporary kubectl bash completion with the following command, run it:

```
user@ubuntu:~$ source <(kubectl completion bash)
user@ubuntu:~$
```

Now try kubectl again:

```
user@ubuntu:~$ kubectl get
apiservices.apiregistration.k8s.io      daemonsets.apps
leases.coordination.k8s.io
...
user@ubuntu:~$ kubectl get
```

That is much better! You can now type "kubectl get ser" and it will autocomplete to "kubectl get service".

In a new shell, list the currently running services, deployments, replica sets, and pods on your cluster:

```
user@ubuntu:~$ kubectl get service,deployments,replicasets,pods

NAME                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/kubernetes  ClusterIP     10.96.0.1     <none>         443/TCP    28m

user@ubuntu:~$
```

The only service running in our cluster is the *kubernetes* service itself. We have no deployments, replica sets, or pods yet (in our namespace). Do the same for the resources under the kube-system namespace, more on namespaces later.

```
user@ubuntu:~$ kubectl get service,deployments,replicaset,pods --namespace=kube-system

NAME                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/kube-dns    ClusterIP     10.96.0.10    <none>         53/UDP,53/TCP,9153/TCP    28m

NAME                READY  UP-TO-DATE  AVAILABLE  AGE
deployment.apps/coredns  2/2    2           2          28m

NAME                DESIRED  CURRENT  READY  AGE
replicaset.apps/coredns-5644d7b6d9  2        2        2      28m

NAME                READY  STATUS    RESTARTS  AGE
pod/coredns-5644d7b6d9-b4rnz  1/1    Running   0         28m
pod/coredns-5644d7b6d9-lxdqv  1/1    Running   0         28m
pod/etcd-ubuntu        1/1    Running   0         27m
pod/kube-apiserver-ubuntu  1/1    Running   0         27m
pod/kube-controller-manager-ubuntu  1/1    Running   0         27m
pod/kube-proxy-npxks      1/1    Running   0         28m
```



```
pod/kube-scheduler-ubuntu      1/1      Running  0          27m
pod/weave-net-rvhvk            2/2      Running  0          12m

user@ubuntu:~$
```

We can view all namespaces via `--all-namespaces` (if we have permission).

To test our cluster lets run a single container pod. When configured with the Docker Engine as the container manager, we can run any container image that Docker has preinstalled or knows how to download.

```
user@ubuntu:~$ kubectl run my-nginx --generator=run-pod/v1 --image=nginx:1.11 --port=80

pod/my-nginx created

user@ubuntu:~$
```

The pod name is "my-nginx" and the image we used is "nginx", an official image pulled from Docker Hub by the Docker Engine in the background. The port switch tells Kubernetes the service port for our pod which will allow us to share the service with its users over that port (the program must actually use that port for this to work).

List the pods running on the cluster:

```
user@ubuntu:~$ kubectl get pods

NAME          READY   STATUS             RESTARTS   AGE
my-nginx      0/1     ContainerCreating   0          9s

user@ubuntu:~$ kubectl get pods

NAME          READY   STATUS    RESTARTS   AGE
my-nginx      1/1     Running   0          23s

user@ubuntu:~$
```

This shows that our pods are deployed and up to date. It may take a bit to pull the Docker images (Ready might be 0).

You can use the `docker container ls` subcommand to display the containers running under the Docker Engine:

```
user@ubuntu:~$ docker container ls --filter "name=nginx"

CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS
PORTS         NAMES
2e25387dfb91   nginx         "nginx -g 'daemon of..." 28 seconds ago Up 27
seconds      k8s_my-nginx_my-nginx_default_2fbe6ccd-f305-4bb0-9d56-
64e5407ee30b_0
1097d5dc9ecf   k8s.gcr.io/pause:3.1  "/pause"                 41 seconds ago Up 40
seconds      k8s_POD_my-nginx_default_2fbe6ccd-f305-4bb0-9d56-
64e5407ee30b_0

user@ubuntu:~$
```

As you can see, while our `run` subcommand requested that Kubernetes run a container but 2 containers were launched at that time.

In Kubernetes, each Pod instance has an infrastructure container, which is the first container that the `kubelet` instantiates. The infrastructure container uses the image "k8s.gcr.io/pause:3.1" and acquires the pod's IP as well as a pod wide network and IPC namespace. All of the other containers in the pod then join the infrastructure container's network (`--net`) and IPC (`--ipc`) namespace allowing containers in the pod to easily communicate. The initial process (`"/pause"`) that runs in the infrastructure container does nothing, its sole purpose is to act as the anchor for the pod and its shared namespaces.

You can learn more about the pause container by looking at the source and ultimately what is "pause()".

- <https://github.com/kubernetes/kubernetes/tree/master/build/pause>
- <https://github.com/kubernetes/kubernetes/blob/master/build/pause/pause.c>

- `man 2 pause` or <http://man7.org/linux/man-pages/man2/pause.2.html>

The Docker listing shows us 2 containers, the pod having an infrastructure container (pause) and the container we asked for (nginx).

Kubernetes gives each pod a name and reports on the pod status, the number of times the pod has been restarted and the pod's uptime. You can find the pod names embedded in the container names displayed by the `docker container ls` command:

```
user@ubuntu:~$ docker container ls --filter "name=nginx" --format "{{.Names}}"

k8s_my-nginx_my-nginx_default_2fbe6ccd-f305-4bb0-9d56-64e5407ee30b_0
k8s_POD_my-nginx_default_2fbe6ccd-f305-4bb0-9d56-64e5407ee30b_0

user@ubuntu:~$
```

Try killing the nginx container using the `docker container kill` subcommand and the ID of the underlying container based on the nginx image.

```
user@ubuntu:~$ docker container kill \
$(docker container ls --filter "ancestor=nginx:1.11" --format {{.ID}} | head -1)

2e25387dfb91

user@ubuntu:~$
```

```
user@ubuntu:~$ docker container ls --filter "name=nginx"

CONTAINER ID        IMAGE               COMMAND                  CREATED              STATUS
PORTS              NAMES
a21bcfc9c4d7        5766334bdaa0       "nginx -g 'daemon of..." 10 seconds ago      Up 9
seconds           k8s_my-nginx_my-nginx_default_2fbe6ccd-f305-4bb0-9d56-
64e5407ee30b_1
1097d5dc9ecf        k8s.gcr.io/pause:3.1 "/pause"                2 minutes ago       Up 2
minutes          k8s_POD_my-nginx_default_2fbe6ccd-f305-4bb0-9d56-
64e5407ee30b_0

user@ubuntu:~$
```

We can tell by the created time we have a new container. If you were fast enough, you may have seen the previous container exited. Docker terminates the container specified but Kubernetes has no knowledge of this action. When the Kubelet process, responsible for the pods assigned to this node, sees the missing container, it simply reruns the nginx image.

After some time, if you run the previous command with the `-a` flag, we can see the previous killed container and the newly created one.

```
user@ubuntu:~$ docker container ls -a --filter "name=nginx"

CONTAINER ID        IMAGE               COMMAND                  CREATED              STATUS
PORTS              NAMES
a21bcfc9c4d7        5766334bdaa0       "nginx -g 'daemon of..." 52 seconds ago      Up 51
seconds           k8s_my-nginx_my-nginx_default_2fbe6ccd-f305-4bb0-
9d56-64e5407ee30b_1
2e25387dfb91        nginx              "nginx -g 'daemon of..." 2 minutes ago       Exited
(137) 52 seconds ago k8s_my-nginx_my-nginx_default_2fbe6ccd-f305-4bb0-
9d56-64e5407ee30b_0
1097d5dc9ecf        k8s.gcr.io/pause:3.1 "/pause"                2 minutes ago       Up 2
minutes          k8s_POD_my-nginx_default_2fbe6ccd-f305-4bb0-9d56-
64e5407ee30b_0

user@ubuntu:~$
```

Notice that we killed container `ae6e77e2b8c2` in the example but the new container `e6ecace9aba4` was created to replace it. Kubernetes does not "resurrect" containers that have failed. This is important because the container's state may be the reason it failed. Rather, Kubernetes runs a fresh copy of the original image, ensuring the container has a clean new internal state (cattle not

pets!).

Having killed the container on the Docker level, check to see how Kubernetes handled the event:

```
user@ubuntu:~$ kubectl get pod

NAME          READY   STATUS    RESTARTS   AGE
my-nginx      1/1     Running   1           3m16s

user@ubuntu:~$
```

When Kubernetes saw that the pod's container had become unavailable, it restarted the pod (and not the container!), incrementing the amount of restarts. The pod itself is still the same, as seen by its AGE not rotating to a smaller number, but it restarted as it launched a new container.

## 4. Create a Service

In modern software engineering terms, a **service** is an encapsulated set of functionality made available to consumers through an API. The problem with our nginx application at present is that when containers die new ones are created. The fact that there are multiple containers and that containers come and go makes using the app difficult.

To simplify things Kubernetes makes it possible for us to expose our pods as a Service. The `kubectl expose` command does this.

Expose the my-nginx pod as a service:

```
user@ubuntu:~$ kubectl expose $(kubectl get pod -o=name) --port=80

service/my-nginx exposed

user@ubuntu:~$
```

This causes Kubernetes to create a conceptual Service for our pods, exposing the set of pods as a single endpoint for users. Use the `get services` subcommand to display your service.

```
user@ubuntu:~$ kubectl get services

NAME          TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes    ClusterIP     10.96.0.1       <none>           443/TCP      33m
my-nginx      ClusterIP     10.106.240.235  <none>           80/TCP       10s

user@ubuntu:~$
```

Kubernetes has given our service a virtual IP (VIP) address and it will now distribute client connections across any running my-nginx pods.

To test the Service try curling it:

```
user@ubuntu:~$ NX_CIP=$(kubectl get services -o=custom-
columns=NAME:.spec.clusterIP,NAME:.metadata.name \
| grep nginx | awk '{print $1}') && echo $NX_CIP

10.106.240.235

user@ubuntu:~$ curl -I $NX_CIP

HTTP/1.1 200 OK
Server: nginx/1.11.13
Date: Wed, 08 Jan 2020 20:49:02 GMT
Content-Type: text/html
Content-Length: 612
Last-Modified: Tue, 04 Apr 2017 15:01:57 GMT
Connection: keep-alive
ETag: "58e3b565-264"
Accept-Ranges: bytes
```

```
user@ubuntu:~$
```

Success!

## 5. Pod exec

While Kubernetes delegates all of the direct container operations to the container manager (usually Docker) it does pass through some useful container features.

For example, imagine you need to discover the distro of one of your pods' containers. You can use the `kubectl exec` subcommand to run arbitrary commands within a pod.

Try listing the running pods and then executing the `cat /etc/os-release` command within one of your pods.

```
user@ubuntu:~$ kubectl get pods

NAME          READY   STATUS    RESTARTS   AGE
my-nginx      1/1     Running   1           4m37s

user@ubuntu:~$ kubectl exec my-nginx -- cat /etc/os-release

PRETTY_NAME="Debian GNU/Linux 8 (jessie)"
NAME="Debian GNU/Linux"
VERSION_ID="8"
VERSION="8 (jessie)"
ID=debian
HOME_URL="http://www.debian.org/"
SUPPORT_URL="http://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"

user@ubuntu:~$
```

Running `cat /etc/os-release` via `kubectl exec` produces the information we needed. The `exec` subcommand chooses the first container within the pod to execute the command. The command to run on the pod was separated from the rest of the `kubectl` invocation with `--`.

If you would like to execute the command within a specific container in a multi-container pod, you can use the `-c` switch.

Try it first find a deployed pod with more than one container:

```
user@ubuntu:~$ kubectl get pods --all-namespaces

NAMESPACE     NAME                                     READY   STATUS    RESTARTS   AGE
default       my-nginx                                1/1     Running   1           5m58s
kube-system   coredns-5644d7b6d9-b4rnz               1/1     Running   0           34m
kube-system   coredns-5644d7b6d9-lxdqv               1/1     Running   0           34m
kube-system   etcd-ubuntu                             1/1     Running   0           33m
kube-system   kube-apiserver-ubuntu                   1/1     Running   0           34m
kube-system   kube-controller-manager-ubuntu          1/1     Running   0           33m
kube-system   kube-proxy-npxks                        1/1     Running   0           34m
kube-system   kube-scheduler-ubuntu                   1/1     Running   0           33m
kube-system   weave-net-rvhvk                         2/2     Running   0           18m

user@ubuntu:~$
```

The weave-net pod for our cluster's networking has two containers in it.

Try to check the os-release on that pod, making sure to use the `--namespace kube-system` so `kubectl` knows which namespace to look:

```
user@ubuntu:~$ kubectl --namespace kube-system exec weave-net-rvhvk -- cat /etc/os-release

Defaulting container name to weave.
Use 'kubectl describe pod/weave-net-rvhvk -n kube-system' to see all of the containers in this
```

```
pod.  
NAME="Alpine Linux"  
ID=alpine  
VERSION_ID=3.8.4  
PRETTY_NAME="Alpine Linux v3.8"  
HOME_URL="http://alpinelinux.org"  
BUG_REPORT_URL="http://bugs.alpinelinux.org"  
  
user@ubuntu:~$
```

Let's do as it suggest and see what containers are inside the weave-net pod. The *describe pod* command will give you a list of the containers within the pod, but we can also use `kubectl get` to retrieve a pod's JSON output.

Use `kubectl get pod` with the `-o json` option to retrieve information about the weave-net pod:

```
user@ubuntu:~$ kubectl --namespace kube-system get pod weave-net-rvhvk -o json  
{  
  "apiVersion": "v1",  
  "kind": "Pod",  
  "metadata": {  
    "creationTimestamp": "2020-01-08T20:31:47Z",  
    "generateName": "weave-net-",  
    "labels": {  
      "controller-revision-hash": "7f54576664",  
      "name": "weave-net",  
      "pod-template-generation": "1"  
    },  
    "name": "weave-net-rvhvk",  
    "namespace": "kube-system",  
    "ownerReferences": [  
      {  
        "apiVersion": "apps/v1",  
        "blockOwnerDeletion": true,  
        "controller": true,  
        "kind": "DaemonSet",  
        "name": "weave-net",  
        "uid": "2fb69d44-878b-4b93-b740-e5252ba94cb5"  
      }  
    ],  
    "resourceVersion": "1627",  
    "selfLink": "/api/v1/namespaces/kube-system/pods/weave-net-rvhvk",  
    "uid": "0028c4b6-26bb-4d4b-91dc-859d2a4272c9"  
  },  
  "spec": {  
    ...  
    "hostIP": "192.168.228.157",  
    "phase": "Running",  
    "podIP": "192.168.228.157",  
    "podIPs": [  
      {  
        "ip": "192.168.228.157"  
      }  
    ],  
    "qosClass": "Burstable",  
    "startTime": "2020-01-08T20:31:47Z"  
  }  
}  
  
user@ubuntu:~$
```

That's a lot of information to sift through, as the entire running pod's spec is presented to you in a single JSON document. You can use a JSON processor like `jq` to filter for specific information from the JSON output.

Use `apt` to install `jq`:

```
user@ubuntu:~$ sudo apt install jq -y
```

```
[sudo] password for user:
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  libonig2
The following NEW packages will be installed:
  jq libonig2
0 upgraded, 2 newly installed, 0 to remove and 262 not upgraded.

...

Setting up jq (1.5+dfsg-1ubuntu0.1) ...
Processing triggers for libc-bin (2.23-0ubuntu3) ...

user@ubuntu:~$
```

And now try to retrieve just the container names from the weave-net pod by piping the JSON output of weave-net to jq:

```
user@ubuntu:~$ kubectl --namespace kube-system get pod weave-net-rvhvk -o json | jq
".spec.containers[].name" -r

weave
weave-npc

user@ubuntu:~$
```

Perfect. You will be using the `-o json` option with `jq` throughout the rest of the labs to gather information about your pods.

Use the `-c` switch to display the `os-release` file from the weave-npc container in the weave-net pod:

```
user@ubuntu:~$ kubectl --namespace kube-system exec -c weave-npc weave-net-rvhvk -- cat /etc/os-release

NAME="Alpine Linux"
ID=alpine
VERSION_ID=3.8.4
PRETTY_NAME="Alpine Linux v3.8"
HOME_URL="http://alpinelinux.org"
BUG_REPORT_URL="http://bugs.alpinelinux.org"

user@ubuntu:~$
```

Being able to execute commands on pods ad-hoc can be very useful in debugging and normal operation scenarios.

## 6. System Logs

Each of the services composing our Kubernetes cluster emits a log file. In the current configuration, the `kubelet` log is controlled by systemd.

You can use the `journalctl` command to tail ( `-n` ) the output for the kubelet service unit ( `-u` )

```
user@ubuntu:~$ journalctl -n 400 --no-pager -u kubelet.service | grep -v "no observation"

-- Logs begin at Wed 2020-01-08 11:49:44 PST, end at Wed 2020-01-08 12:56:36 PST. --
Jan 08 12:15:59 ubuntu kubelet[68925]: W0108 12:15:59.068301 68925 cni.go:237] Unable to
update cni config: no networks found in /etc/cni/net.d
Jan 08 12:15:59 ubuntu kubelet[68925]: E0108 12:15:59.655595 68925 kubelet.go:2187] Container
runtime network not ready: NetworkReady=false reason:NetworkPluginNotReady message:docker:
network plugin is not ready: cni config uninitialized
Jan 08 12:16:04 ubuntu kubelet[68925]: W0108 12:16:04.068919 68925 cni.go:237] Unable to
update cni config: no networks found in /etc/cni/net.d
Jan 08 12:16:04 ubuntu kubelet[68925]: E0108 12:16:04.663617 68925 kubelet.go:2187] Container
runtime network not ready: NetworkReady=false reason:NetworkPluginNotReady message:docker:
```

```
network plugin is not ready: cni config uninitialized
Jan 08 12:16:09 ubuntu kubelet[68925]: W0108 12:16:09.069770 68925 cni.go:237] Unable to
update cni config: no networks found in /etc/cni/net.d

...

user@ubuntu:~$
```

The rest of our services are running as containers.

We can use the `kubectl logs` command to display log output from our pods. Remember that Kubernetes system services run within the `kube-system` namespace by convention.

List the pods in the `kube-system` namespace:

```
user@ubuntu:~$ kubectl get pods --namespace=kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-5644d7b6d9-b4rnz	1/1	Running	0	43m
coredns-5644d7b6d9-lxdqv	1/1	Running	0	43m
etcd-ubuntu	1/1	Running	0	42m
kube-apiserver-ubuntu	1/1	Running	0	42m
kube-controller-manager-ubuntu	1/1	Running	0	41m
kube-proxy-npxks	1/1	Running	0	43m
kube-scheduler-ubuntu	1/1	Running	0	42m
weave-net-rvhvk	2/2	Running	0	27m

```
user@ubuntu:~$
```

Now display the last 10 lines from the API service:

```
user@ubuntu:~$ kubectl logs --namespace=kube-system --tail=10 kube-apiserver-ubuntu
```

```
I0108 20:15:30.509997 1 storage_scheduling.go:148] all system priority classes are created
successfully or already exist.
I0108 20:15:30.771391 1 controller.go:606] quota admission added evaluator for:
roles.rbac.authorization.k8s.io
I0108 20:15:30.801684 1 controller.go:606] quota admission added evaluator for:
rolebindings.rbac.authorization.k8s.io
W0108 20:15:30.937485 1 lease.go:222] Resetting endpoints for master service "kubernetes"
to [192.168.228.157]
I0108 20:15:30.940434 1 controller.go:606] quota admission added evaluator for: endpoints
I0108 20:15:31.955961 1 controller.go:606] quota admission added evaluator for:
serviceaccounts
I0108 20:15:31.968713 1 controller.go:606] quota admission added evaluator for:
deployments.apps
I0108 20:15:32.310219 1 controller.go:606] quota admission added evaluator for:
daemonsets.apps
I0108 20:15:39.903929 1 controller.go:606] quota admission added evaluator for:
replicasets.apps
I0108 20:15:39.919553 1 controller.go:606] quota admission added evaluator for:
controllerrevisions.apps

user@ubuntu:~$
```

Each Kubernetes service has its own log verbosity and each can be tuned. You can learn much by tracking the operations involved in starting a deployment.

Create a new single pod deployment with a descriptive name can be tracked for activity in the logs:

```
user@ubuntu:~$ kubectl run --generator=run-pod/v1 mylogtracker --image nginx:1.11

pod/mylogtracker created

user@ubuntu:~$
```

Again list the k8s system services, from here we can pick which logs to search for our new pod.

```
user@ubuntu:~$ kubectl get pod --namespace=kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-5644d7b6d9-b4rnz	1/1	Running	0	43m
coredns-5644d7b6d9-lxdqv	1/1	Running	0	43m
etcd-ubuntu	1/1	Running	0	42m
kube-apiserver-ubuntu	1/1	Running	0	43m
kube-controller-manager-ubuntu	1/1	Running	0	42m
kube-proxy-npxks	1/1	Running	0	43m
kube-scheduler-ubuntu	1/1	Running	0	42m
weave-net-rvhvk	2/2	Running	0	27m

```
user@ubuntu:~$
```

Try the controller manager server first:

```
user@ubuntu:~$ kubectl logs --namespace=kube-system kube-controller-manager-ubuntu | grep mylogtracker
```

```
user@ubuntu:~$
```

Controller manager only deals with replicated pods (ones using a controller); there won't be anything here for us.

Now take a look at the kubelet log:

```
user@ubuntu:~$ journalctl -u kubelet.service | grep mylogtracker
```

```
Jan 08 12:59:27 ubuntu kubelet[68925]: I0108 12:59:27.649965 68925 reconciler.go:207] operationExecutor.VerifyControllerAttachedVolume started for volume "default-token-7bqf5" (UniqueName: "kubernetes.io/secret/306f090b-46fc-4512-8f4a-a8b18dcd1a9e-default-token-7bqf5") pod "mylogtracker" (UID: "306f090b-46fc-4512-8f4a-a8b18dcd1a9e")
```

```
user@ubuntu:~$
```

Kubelet only reports information about our pod's secret, which is mounted as volume.

You can also view the events taking place within the Kubernetes cluster itself using the events resource type.

Try getting events with `kubectl` :

```
user@ubuntu:~$ kubectl get events --sort-by='{.lastTimestamp}'
```

LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
44m	Normal	NodeHasSufficientPID	node/ubuntu	Node ubuntu status is now:
44m	Normal	NodeHasSufficientPID	node/ubuntu	Node ubuntu status is now:
44m	Normal	NodeHasNoDiskPressure	node/ubuntu	Node ubuntu status is now:
44m	Normal	NodeHasNoDiskPressure	node/ubuntu	Node ubuntu status is now:
44m	Normal	NodeHasSufficientMemory	node/ubuntu	Node ubuntu status is now:
44m	Normal	NodeHasSufficientMemory	node/ubuntu	Node ubuntu status is now:
44m	Normal	RegisteredNode	node/ubuntu	Node ubuntu event: Registered
44m	Normal	Node ubuntu in Controller	node/ubuntu	Starting kube-proxy.
44m	Normal	Starting	node/ubuntu	Starting kube-proxy.
28m	Normal	NodeReady	node/ubuntu	Node ubuntu status is now:
15m	Normal	NodeReady	node/ubuntu	Node ubuntu status is now:
15m	Normal	Scheduled	pod/my-nginx	Successfully assigned
15m	Normal	default/my-nginx to ubuntu	pod/my-nginx	Successfully assigned
15m	Normal	Pulling	pod/my-nginx	Pulling image "nginx:1.11"
15m	Normal	Pulled	pod/my-nginx	Successfully pulled image
13m	Normal	"nginx:1.11"	pod/my-nginx	Successfully pulled image
13m	Normal	Pulled	pod/my-nginx	Container image "nginx:1.11"
13m	Normal	already present on machine	pod/my-nginx	Container image "nginx:1.11"
13m	Normal	Started	pod/my-nginx	Started container my-nginx
13m	Normal	Created	pod/my-nginx	Created container my-nginx



50s	Normal	Scheduled	pod/mylogtracker	Successfully assigned
default/mylogtracker		to ubuntu		
49s	Normal	Created	pod/mylogtracker	Created container mylogtracker
49s	Normal	Pulled	pod/mylogtracker	Container image "nginx:1.11"
already present on machine				
48s	Normal	Started	pod/mylogtracker	Started container mylogtracker

```
user@ubuntu:~$
```

While your events will be different you can see the value of the cluster event log. You can display event data associated with a given resource by supplying its name. You can also control the output format.

For example to make the data machine readable you could output it in JSON:

```
user@ubuntu:~$ kubectl get events -o json | grep mylogtracker -A5 | tail

      "creationTimestamp": "2020-01-08T20:59:29Z",
      "name": "mylogtracker.15e805007f603ffb",
      "namespace": "default",
      "resourceVersion": "3674",
      "selfLink": "/api/v1/namespaces/default/events/mylogtracker.15e805007f603ffb",
      "uid": "e6afaa59-9294-4f24-b9f2-6340dd647fd0"
    },
    "reason": "Started",
    "reportingComponent": "",
    "reportingInstance": "",
  }
}
```

```
user@ubuntu:~$
```

## 7. Cleaning Up

Now that we have given our new cluster a good test we can clean up by deleting the service and deployments we have created. The `kubectl delete` subcommand allows you to delete objects you have created in the cluster.

To begin, delete the *my-nginx* Service:

```
user@ubuntu:~$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	45m
my-nginx	ClusterIP	10.106.240.235	<none>	80/TCP	12m

```
user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl delete service my-nginx

service "my-nginx" deleted

user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	46m

```
user@ubuntu:~$
```

Do not delete the kubernetes service.

Next we can delete the pods:

```
user@ubuntu:~$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
my-nginx	1/1	Running	1	17m
mylogtracker	1/1	Running	0	2m31s

```
user@ubuntu:~$
```

You can specify multiple resources to by placing spaces between each resource:

```
user@ubuntu:~$ kubectl delete pod my-nginx mylogtracker
```

pod "my-nginx" deleted  
pod "mylogtracker" deleted

```
user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl get pods
```

No resources found in default namespace.

```
user@ubuntu:~$
```

You Kubernetes cluster should now be cleaned up and ready for the next lab:

```
user@ubuntu:~$ kubectl get services,deployments,replicasets,pods
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	47m

```
user@ubuntu:~$
```

Be sure to leave the `service/kubernetes` service!

Congratulations, you have completed the lab!

*Copyright (c) 2013-2020 RX-M LLC, Cloud Native Consulting, all rights reserved*