

Kubernetes

Lab 4 – Deployments and Replica Sets

In this lab we will explore the nature of Kubernetes deployments and replica sets and how to work with them.

Deployments

A deployment provides declarative updates for pods and replica sets. You describe the desired state in a deployment object, and the deployment controller will change the actual state to the desired state at a controlled rate for you. You can define deployments to create new resources, or replace existing ones by new ones. Typical uses:

- bring up a replica set and (indirectly) its pods
- capturing the results and status of a deployment
- updating an existing deployment to recreate pods with a new image (rolling updates)
- rolling back to an earlier deployment revision if the current deployment isn't stable
- pausing and resuming a deployment

ReplicaSets

Replica sets (RS) supersede the older replication controller (RC) resource type. Replica sets support the set-based selectors as well as equality-based selector requirements (RCs only supported equality.) While replica sets can be used independently, they are mainly used by deployments as a mechanism to orchestrate pod creation, deletion, and updates. When you use deployments you don't have to worry about managing the replica sets that they create; deployments own and manage their replica sets.

ReplicaSets ensure that a specified number of pod "replicas" are running at all times. If there are too many, it will kill some. If there are too few, it will start more. Unlike in the case where a user directly created pods, a ReplicaSet replaces pods that are deleted or terminated for any reason, such as in the case of node failure or disruptive node maintenance (e.g. a kernel upgrade, etc.)

For this reason the Kubernetes team recommends that you use a Deployment/ReplicaSet even if your application requires only a single pod. ReplicaSets are like process supervisors in many ways but monitor processes on multiple nodes at once. A ReplicaSet delegates local container restarts to some agent on the node (e.g., Kubelet or Docker.)

A ReplicaSet is only appropriate for pods with *RestartPolicy = Always* (if the RestartPolicy is not set, the default value is *Always*.) A ReplicaSet will refuse to instantiate any pod that has a different restart policy.

A ReplicaSet will never terminate on its own, but it isn't expected to be as long-lived as services. Services may be composed of pods controlled by multiple ReplicaSets, and it is expected that many ReplicaSets may be created and destroyed over the lifetime of a service (for instance, to perform an update of pods that run the service.) Both services themselves and their clients should remain oblivious to the ReplicaSets that maintain the pods of the services.

Now to create some Deployments/ReplicaSets.

1. A Simple Deployment

As a first exploratory step let's create a simple deployment which stands up three nginx pods. Create a config file similar to the following to accomplish this task:

```
user@ubuntu:~$ cd ~

user@ubuntu:~$ mkdir ~/dep && cd ~/dep

user@ubuntu:~/dep$ nano mydep.yaml && cat mydep.yaml

apiVersion: apps/v1
kind: Deployment
metadata:
  name: website
  labels:
    bu: sales
spec:
  replicas: 3
```

```

selector:
  matchLabels:
    appname: webserver
    targetenv: demo
template:
  metadata:
    labels:
      appname: webserver
      targetenv: demo
  spec:
    containers:
    - name: podweb
      image: nginx:1.7.9
      ports:
      - containerPort: 80

```

```
user@ubuntu:~/dep$
```

Deployments were promoted to the apps/v1 API in K8s 1.9 but were added to Kubernetes 1.2 and are the go forward solution for deploying replicated pods. The spec for Replication Controllers (part of the v1 API) is almost the same as the spec for Deployments though deployments add a few key features such as the ability to specify upgrades declaratively. The specification for Deployments can be found here:

<https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.15/#deployment-v1-apps>

Now create the Deployment using the `kubectl apply` subcommand and verify that the Deployment, its ReplicaSet and pods are up with the `get` subcommand:

```

user@ubuntu:~/dep$ kubectl apply -f mydep.yaml

deployment.apps/website created

user@ubuntu:~/dep$

```

List the deployments, replicaset and pods. You can abbreviate certain API objects when calling them, as shown below:

```

user@ubuntu:~/dep$ kubectl get deploy,rs,po

```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/website	3/3	3	3	15s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/website-5577f87457	3	3	3	15s

NAME	READY	STATUS	RESTARTS	AGE
pod/website-5577f87457-6gxrf	1/1	Running	0	15s
pod/website-5577f87457-fswcn	1/1	Running	0	15s
pod/website-5577f87457-vv96d	1/1	Running	0	15s

```

user@ubuntu:~/dep$

```

While everything appears to be running we can verify that there are no scheduling cycles or fail/restart activities by examining the system events. We have viewed resource specific events in the past using the `kubectl describe` subcommand. This time we'll use the `kubectl get events` subcommand to view cluster wide events:

```

user@ubuntu:~/dep$ kubectl get events --sort-by='{metadata.creationTimestamp}' | grep website

```

3m16s	Normal	SuccessfulCreate	replicaset/website-5577f87457	Created pod: website-5577f87457-fswcn
<unknown>	Normal	Scheduled	pod/website-5577f87457-6gxrf	Successfully assigned default/website-5577f87457-6gxrf to ubuntu
3m16s	Normal	SuccessfulCreate	replicaset/website-5577f87457	Created pod: website-5577f87457-vv96d
3m16s	Normal	SuccessfulCreate	replicaset/website-5577f87457	Created pod: website-5577f87457-6gxrf
<unknown>	Normal	Scheduled	pod/website-5577f87457-vv96d	Successfully

```

assigned default/website-5577f87457-vv96d to ubuntu
<unknown> Normal Scheduled pod/website-5577f87457-fswcn Successfully
assigned default/website-5577f87457-fswcn to ubuntu
3m16s Normal ScalingReplicaSet deployment/website Scaled up
replica set website-5577f87457 to 3
3m14s Normal Pulling pod/website-5577f87457-fswcn Pulling image
"nginx:1.7.9"
3m14s Normal Pulling pod/website-5577f87457-6gxrf Pulling image
"nginx:1.7.9"
3m14s Normal Pulling pod/website-5577f87457-vv96d Pulling image
"nginx:1.7.9"
3m5s Normal Pulled pod/website-5577f87457-6gxrf Successfully
pulled image "nginx:1.7.9"
3m5s Normal Created pod/website-5577f87457-6gxrf Created
container podweb
3m5s Normal Started pod/website-5577f87457-6gxrf Started
container podweb
3m4s Normal Created pod/website-5577f87457-vv96d Created
container podweb
3m4s Normal Pulled pod/website-5577f87457-vv96d Successfully
pulled image "nginx:1.7.9"
3m3s Normal Started pod/website-5577f87457-vv96d Started
container podweb
3m2s Normal Started pod/website-5577f87457-fswcn Started
container podweb
3m2s Normal Created pod/website-5577f87457-fswcn Created
container podweb
3m2s Normal Pulled pod/website-5577f87457-fswcn Successfully
pulled image "nginx:1.7.9"

user@ubuntu:~/dep$

```

Checking the event log occasionally will help you identify normal cluster patterns and make it possible for you to spot anomalies more easily when debugging.

The replica set began with a scale of 3, causing 3 instances of the pod template to get scheduled. Replica sets ensure that some number of instances of the pod template are always running. Try deleting a pod (use the pod name displayed by `kubectl get pods`).

```

user@ubuntu:~/dep$ kubectl delete $(kubectl get pod -o name |tail -1)

pod "website-5577f87457-vv96d" deleted

user@ubuntu:~/dep$ kubectl get pods

NAME                                READY   STATUS    RESTARTS   AGE
website-5577f87457-6gxrf            1/1     Running   0           4m20s
website-5577f87457-fswcn            1/1     Running   0           4m20s
website-5577f87457-mb7cv            1/1     Running   0           20s

user@ubuntu:~/dep$

```

You might ask: "why would Kubernetes let someone delete the pod if it will just restart it?". There are many reasons you might want to delete a pod:

- The pod ran into problems and you want to generate a new replacement.
- The current node has problems and you want Kubernetes to reschedule this particular pod somewhere else.
- A change was made to another Kubernetes resource the pod relied on, like a configMap or volume mount, which require a pod to be recreated to take effect in many cases.

To actually terminate our pod permanently we must delete the deployment, the deployment controls the replica set, the replica set controls the pods.

When many resources are running on a cluster it can be advantageous to restrict output to a certain set of resources. The Kubernetes labeling system makes this easy. The `-l` switch can be used with the `kubectl get` subcommand to filter output by label.

Try listing all pods with the `--show-labels` options to show the pod labels:

```
user@ubuntu:~/dep$ kubectl get pods --show-labels --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	LABELS
default	website-5577f87457-6gxrf	1/1	Running	0	4m40s	
appname=webserver,pod-template-hash=5577f87457,targetenv=demo						
default	website-5577f87457-fswcn	1/1	Running	0	4m40s	
appname=webserver,pod-template-hash=5577f87457,targetenv=demo						
default	website-5577f87457-mb7cv	1/1	Running	0	40s	
appname=webserver,pod-template-hash=5577f87457,targetenv=demo						
kube-system	coredns-5644d7b6d9-b4rnz	1/1	Running	1	80m	k8s-
app=kube-dns,pod-template-hash=5644d7b6d9						
kube-system	coredns-5644d7b6d9-lxdqv	1/1	Running	1	80m	k8s-
app=kube-dns,pod-template-hash=5644d7b6d9						
kube-system	etcd-ubuntu	1/1	Running	1	79m	
component=etcd,tier=control-plane						
kube-system	kube-apiserver-ubuntu	1/1	Running	1	79m	
component=kube-apiserver,tier=control-plane						
kube-system	kube-controller-manager-ubuntu	1/1	Running	1	79m	
component=kube-controller-manager,tier=control-plane						
kube-system	kube-proxy-npxks	1/1	Running	1	80m	controller-
revision-hash=57cd7f4c65,k8s-app=kube-proxy,pod-template-generation=1						
kube-system	kube-scheduler-ubuntu	1/1	Running	1	79m	
component=kube-scheduler,tier=control-plane						
kube-system	weave-net-rvhvk	2/2	Running	3	64m	controller-
revision-hash=7f54576664,name=weave-net,pod-template-generation=1						

```
user@ubuntu:~/dep$
```

Now try filtering by the "appname" label key we assigned to all of our pods in the pod template metadata, searching all namespaces:

```
user@ubuntu:~/dep$ kubectl get pods --show-labels --all-namespaces -l appname
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	LABELS
default	website-5577f87457-6gxrf	1/1	Running	0	5m11s	
appname=webserver,pod-template-hash=5577f87457,targetenv=demo						
default	website-5577f87457-fswcn	1/1	Running	0	5m11s	
appname=webserver,pod-template-hash=5577f87457,targetenv=demo						
default	website-5577f87457-mb7cv	1/1	Running	0	71s	
appname=webserver,pod-template-hash=5577f87457,targetenv=demo						

```
user@ubuntu:~/dep$
```

You can also filter by key and value:

```
user@ubuntu:~/dep$ kubectl get pods --all-namespaces -l targetenv=demo
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	website-5577f87457-6gxrf	1/1	Running	0	5m35s
default	website-5577f87457-fswcn	1/1	Running	0	5m35s
default	website-5577f87457-mb7cv	1/1	Running	0	95s

```
user@ubuntu:~/dep$
```

You can filter by pod name:

```
user@ubuntu:~/dep$ kubectl get $(kubectl get pods -o name | head -1)
```

NAME	READY	STATUS	RESTARTS	AGE
website-5577f87457-6gxrf	1/1	Running	0	5m58s

```
user@ubuntu:~/dep$
```

Our pod has labels we have added and the Kubernetes infrastructure may add labels as well:

```
user@ubuntu:~/dep$ kubectl describe $(kubectl get pods -o name | head -1) | grep -A2 -i label

Labels:      appname=webserver
             pod-template-hash=5577f87457
             targetenv=demo

user@ubuntu:~/dep$
```

Unfortunately describe doesn't allow for JSON output. Good news, though, `get` does.

```
user@ubuntu:~/dep$ kubectl get $(kubectl get pods -o name | head -1) -o json | jq
.metadata.labels
```

```
{
  "appname": "webserver",
  "pod-template-hash": "5577f87457",
  "targetenv": "demo"
}
```

```
user@ubuntu:~/dep$
```

- Why do each of the filters above work or not work?
- Enter a command to display all of the pods with either the "demo" or "prod" value for targetenv
- Find all pods other than those with the "demo" or "prod" value for targetenv
- Enter a command to display all of the pods with either the "demo" or "prod" value for targetenv and the appname key set to webserver

2. Checking status of a Deployment

We have seen previously how to check the status of a deployment.

```
user@ubuntu:~/dep$ kubectl get deploy

NAME      READY   UP-TO-DATE   AVAILABLE   AGE
website   3/3     3            3           6m43s

user@ubuntu:~/dep$
```

Now we take an slightly more application-centric view.

```
user@ubuntu:~/dep$ kubectl rollout status deploy/website

deployment "website" successfully rolled out

user@ubuntu:~/dep$
```

Rollouts are used to update a given set of Pods, the ones controlled by this Deployment's replica set. It reports success when all the currently deployed Pods match what is expected in the current deployment. In k8s technical terms these conditions are all true:

- `.status.observedGeneration >= .metadata.generation`
- `.status.updatedReplicas == .spec.replicas`
- `.spec.availableReplicas >= minimum required`

3. Updating a Deployment

We are using nginx 1.7.9 in our example, lets update to 1.9.1.

```

user@ubuntu:~/dep$ kubectl set image deploy/website podweb=nginx:1.9.1 --record
deployment.apps/website image updated

user@ubuntu:~/dep$

```

Alternative is to use `kubectl edit deployment/website`

Check the status of the rollout (if you're not fast you may not see these updates):

```

user@ubuntu:~/dep$ kubectl rollout status deploy/website

Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "website" rollout to finish: 1 old replicas are pending termination...
deployment "website" successfully rolled out

user@ubuntu:~/dep$

```

```

user@ubuntu:~/dep$ kubectl get deploy/website

NAME      READY   UP-TO-DATE   AVAILABLE   AGE
website   3/3     3            3           7m29s

user@ubuntu:~/dep$

```

Look at the Replica Sets & Pods:

```

user@ubuntu:~/dep$ kubectl get rs,pod

NAME                                     DESIRED   CURRENT   READY   AGE
replicaset.apps/website-5577f87457      0         0         0       7m44s
replicaset.apps/website-769bf6f999      3         3         3       36s

NAME                                     READY   STATUS    RESTARTS   AGE
pod/website-769bf6f999-6m54n            1/1     Running   0          24s
pod/website-769bf6f999-78nlg            1/1     Running   0          36s
pod/website-769bf6f999-rjt2g            1/1     Running   0          22s

user@ubuntu:~/dep$

```

All the pods bear different suffixes, and should be seconds old at this point. When a rollout is performed, the original replicaSet controlling the pods are scaled down to zero, and a new replicaSet bearing the new spec (in this case, a new container image tag) is created and scaled up to replace it.

By describing the deployment we can inspect the events that occurred during the rollout:

```

user@ubuntu:~/dep$ kubectl describe deploy/website | grep -A 10 Events

Events:
  Type     Reason             Age   From                  Message
  ----     -
  Normal   ScalingReplicaSet   8m    deployment-controller  Scaled up replica set website-5577f87457 to 3
  Normal   ScalingReplicaSet   52s   deployment-controller  Scaled up replica set website-769bf6f999 to 1
  Normal   ScalingReplicaSet   40s   deployment-controller  Scaled down replica set website-5577f87457 to 2

```

```

Normal ScalingReplicaSet 40s deployment-controller Scaled up replica set website-
769bf6f999 to 2
Normal ScalingReplicaSet 38s deployment-controller Scaled down replica set website-
5577f87457 to 1
Normal ScalingReplicaSet 38s deployment-controller Scaled up replica set website-
769bf6f999 to 3
Normal ScalingReplicaSet 36s deployment-controller Scaled down replica set website-
5577f87457 to 0

user@ubuntu:~/dep$

```

Note that the rollout was a smooth transition from one set of Pods controlled by our original ReplicaSet to our second set of Pods controlled by the RS.

4. Manually rolling back a deployment

Lets manually revert back to nginx 1.7.9 and check the status:

```

user@ubuntu:~/dep$ kubectl set image deploy/website podweb=nginx:1.7.9 --record

deployment.apps/website image updated

user@ubuntu:~/dep$

```

```

user@ubuntu:~/dep$ kubectl rollout status deploy/website

...

Waiting for rollout to finish: 1 old replicas are pending termination...
Waiting for rollout to finish: 1 old replicas are pending termination...
deployment "website" successfully rolled out

user@ubuntu:~/dep$

```

```

user@ubuntu:~/dep$ kubectl get rs

NAME                DESIRED   CURRENT   READY   AGE
website-5577f87457   3         3         3       9m6s
website-769bf6f999   0         0         0       11

user@ubuntu:~/dep$

```

Notice which RS hash in the pod name is being used.

```

user@ubuntu:~/dep$ kubectl get pods

NAME                                READY   STATUS    RESTARTS   AGE
website-5577f87457-4585h            1/1     Running   0           55s
website-5577f87457-vjf2m            1/1     Running   0           57s
website-5577f87457-vwpcs            1/1     Running   0           5

user@ubuntu:~/dep$

```

Confirm your observations once again in the event log.

```

user@ubuntu:~/dep$ kubectl describe deploy/website | grep -A 15 Events

Events:
  Type    Reason             Age          From              Message
  ----    -
  Normal  ScalingReplicaSet  2m29s       deployment-controller Scaled up replica set

```

```

website-769bf6f999 to 1
Normal ScalingReplicaSet 2m17s deployment-controller Scaled down replica set
website-5577f87457 to 2
Normal ScalingReplicaSet 2m17s deployment-controller Scaled up replica set
website-769bf6f999 to 2
Normal ScalingReplicaSet 2m15s deployment-controller Scaled down replica set
website-5577f87457 to 1
Normal ScalingReplicaSet 2m15s deployment-controller Scaled up replica set
website-769bf6f999 to 3
Normal ScalingReplicaSet 2m13s deployment-controller Scaled down replica set
website-5577f87457 to 0
Normal ScalingReplicaSet 73s deployment-controller Scaled up replica set
website-5577f87457 to 1
Normal ScalingReplicaSet 72s deployment-controller Scaled down replica set
website-769bf6f999 to 2
Normal ScalingReplicaSet 70s (x2 over 9m37s) deployment-controller Scaled up replica set
website-5577f87457 to 3
Normal ScalingReplicaSet 69s (x3 over 72s) deployment-controller (combined from similar
events): Scaled down replica set website-769bf6f999 to 0

user@ubuntu:~/dep$

```

5. Checking rollout history of a Deployment

We can use the *rollout history* subcommand to see what we have been doing to trigger these rollouts:

```

user@ubuntu:~/dep$ kubectl rollout history deploy/website

deployment.apps/website
REVISION  CHANGE-CAUSE
2         kubectl set image deploy/website podweb=nginx:1.9.1 --record=true
3         kubectl set image deploy/website podweb=nginx:1.7.9 --record=true

user@ubuntu:~/dep$

```

Take a detailed look at a previous deployment version:

```

user@ubuntu:~/dep$ kubectl rollout history deploy/website --revision=2

deployment.apps/website with revision #2
Pod Template:
  Labels:      appname=webserver
              pod-template-hash=769bf6f999
              targetenv=demo
  Annotations: kubernetes.io/change-cause: kubectl set image deploy/website podweb=nginx:1.9.1
              --record=true
  Containers:
    podweb:
      Image:      nginx:1.9.1
      Port:       80/TCP
      Host Port:  0/TCP
      Environment:    <none>
      Mounts:         <none>
      Volumes:         <none>

user@ubuntu:~/dep$

```

6. Rolling back to a previous Deployment

Confirm the current version of a container is 1.7.9.

```

user@ubuntu:~/dep$ kubectl get pods -o json | jq .items[0].spec.containers[0].image -r

nginx:1.7.9

```



```
user@ubuntu:~/dep$
```

Revert to previous version/revision, tracking the rollout status by chaining the `kubectl rollout status` command:

```
user@ubuntu:~/dep$ kubectl rollout undo deploy/website && kubectl rollout status deploy/website
deployment.apps/website rolled back
Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "website" rollout to finish: 1 old replicas are pending termination...
deployment "website" successfully rolled out

user@ubuntu:~/dep$
```

Alternative to above is `kubectl rollout undo deployment/website --to-revision=2`

```
user@ubuntu:~/dep$ kubectl get deploy/website

NAME      READY   UP-TO-DATE   AVAILABLE   AGE
website   3/3     3            3           9m49s

user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl describe deploy/website | grep -A 15 Events

Events:
  Type     Reason             Age           From              Message
  ----     -
  Normal   ScalingReplicaSet  4m17s        deployment-controller Scaled down replica set
website-5577f87457 to 2
  Normal   ScalingReplicaSet  4m15s        deployment-controller Scaled down replica set
website-5577f87457 to 1
  Normal   ScalingReplicaSet  4m13s        deployment-controller Scaled down replica set
website-5577f87457 to 0
  Normal   ScalingReplicaSet  3m13s        deployment-controller Scaled up replica set
website-5577f87457 to 1
  Normal   ScalingReplicaSet  3m12s        deployment-controller Scaled down replica set
website-769bf6f999 to 2
  Normal   ScalingReplicaSet  3m10s (x2 over 11m) deployment-controller Scaled up replica set
website-5577f87457 to 3
  Normal   ScalingReplicaSet  43s (x2 over 4m29s) deployment-controller Scaled up replica set
website-769bf6f999 to 1
  Normal   ScalingReplicaSet  41s (x2 over 4m17s) deployment-controller Scaled up replica set
website-769bf6f999 to 2
  Normal   ScalingReplicaSet  39s (x2 over 4m15s) deployment-controller Scaled up replica set
website-769bf6f999 to 3
  Normal   ScalingReplicaSet  37s (x6 over 3m12s) deployment-controller (combined from similar
events): Scaled down replica set website-5577f87457 to 0

user@ubuntu:~/dep$
```

Confirm the container image version has been reverted to 1.9.1:

```
user@ubuntu:~/dep$ kubectl get pods -o json | jq .items[0].spec.containers[0].image -r
nginx:1.9.1

user@ubuntu:~/dep$
```

7. Pausing and resuming a Deployment

In a larger installation, we may be deploying dozens of pods. For our small test it is hard to pause in time, so we chain the commands to hopefully catch it in the act.

```
user@ubuntu:~/dep$ kubectl set image deploy/website podweb=nginx:1.7.9; kubectl rollout pause deploy/website

deployment.apps/website image updated
deployment.apps/website paused

user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl get rs

NAME                DESIRED   CURRENT   READY   AGE
website-5577f87457   1         1         1       12m
website-769bf6f999   3         3         3       5m15s

user@ubuntu:~/dep$
```

It worked! You can see that the original 1.7.9 replicaSet and the 1.9.1 replicaSets both have pods running, despite running the `kubectl set` command.

Check the deployment status with `kubectl rollout`. Since `kubectl rollout` will watch for the change, you will need to back out of the command with CTRL C:

```
user@ubuntu:~/dep$ kubectl rollout status deploy/website

Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have been updated...

^C

user@ubuntu:~/dep$
```

To resume the rollout, you can use the `resume` subcommand for `kubectl rollout`:

```
user@ubuntu:~/dep$ kubectl rollout resume deploy/website && kubectl rollout status deploy/website

deployment.apps/website resumed
Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have been updated...
Waiting for deployment "website" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "website" rollout to finish: 1 old replicas are pending termination...
deployment "website" successfully rolled out

user@ubuntu:~/dep$
```

Now check the replicaSets to see if they have successfully replaced each other:

```
user@ubuntu:~/dep$ kubectl get rs

NAME                DESIRED   CURRENT   READY   AGE
website-5577f87457   3         3         3       13m
website-769bf6f999   0         0         0       6m7s
```

```
user@ubuntu:~/dep$
```

There are no more nginx pods running version 1.9.1, only 1.7.6.

Delete your deployment:

```
user@ubuntu:~/dep$ kubectl delete deploy website  
deployment.apps "website" deleted  
user@ubuntu:~/dep$
```

8. Health Checks

In this step we will create a pod with a health check. Enter and run the following config (*hc.yaml*):

```
user@ubuntu:~/dep$ mkdir ~/hc && cd ~/hc  
user@ubuntu:~/hc$ nano hc.yaml && cat hc.yaml  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: nginx  
  labels:  
    name: nginx  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      name: nginx  
  template:  
    metadata:  
      labels:  
        name: nginx  
    spec:  
      containers:  
      - name: nginx  
        image: nginx:latest  
        ports:  
        - containerPort: 80  
        livenessProbe: # An HTTP health check  
          httpGet:  
            path: /  
            port: 80  
            initialDelaySeconds: 30  
            timeoutSeconds: 1  
user@ubuntu:~/hc$
```

Now run the deployment:

```
user@ubuntu:~/hc$ kubectl apply -f hc.yaml  
deployment.apps/nginx created  
user@ubuntu:~/hc$
```

View your deployment:

```
user@ubuntu:~/hc$ kubectl get deploy,rs,pods  
  
NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
```

```
deployment.apps/nginx 3/3 3 3 13s
```

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/nginx-8cdb45866	3	3	3	13s

NAME	READY	STATUS	RESTARTS	AGE
pod/nginx-8cdb45866-9p8ws	1/1	Running	0	13s
pod/nginx-8cdb45866-lvbkl	1/1	Running	0	13s
pod/nginx-8cdb45866-r89pv	1/1	Running	0	13s

```
user@ubuntu:~/hc$
```

Note that our nginx service listens on port 80 and responds normally to requests for "/", so our health check is passing.

To trigger the health check repair logic, we need to simulate an error condition. By forcing nginx to report a 404, the *httpGet* livenessProbe will fail. We can do this by deleting the nginx configuration file in the nginx container.

Display the events for the first pod in the set:

```
user@ubuntu:~/hc$ kubectl get events --sort-by='{metadata.creationTimestamp}' \
| grep $(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}')
```

32s	Normal	SuccessfulCreate	replicaset/nginx-8cdb45866	Created pod: nginx-8cdb45866-9p8ws
<unknown>	Normal	Scheduled	pod/nginx-8cdb45866-9p8ws	Successfully assigned default/nginx-8cdb45866-9p8ws to ubuntu
29s	Normal	Pulling	pod/nginx-8cdb45866-9p8ws	Pulling image "nginx:latest"
23s	Normal	Started	pod/nginx-8cdb45866-9p8ws	Started container nginx
23s	Normal	Created	pod/nginx-8cdb45866-9p8ws	Created container nginx
23s	Normal	Pulled	pod/nginx-8cdb45866-9p8ws	Successfully pulled image "nginx:latest"

```
user@ubuntu:~/hc$
```

The status is good.

Now lets tell the nginx in the first pod to stop serving the root IRI by deleting the nginx default config:

```
user@ubuntu:~/hc$ kubectl exec -it $(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}') \
-- sh -c "rm /etc/nginx/conf.d/default.conf && nginx -s reload"
```

```
2020/01/08 21:46:40 [notice] 15#15: signal process started
```

```
user@ubuntu:~/hc$
```

Now redisplay the events for the pod:

```
user@ubuntu:~/hc$ kubectl get events --sort-by='{metadata.creationTimestamp}' \
| grep $(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}')
```

66s	Normal	SuccessfulCreate	replicaset/nginx-8cdb45866	Created pod: nginx-8cdb45866-9p8ws
<unknown>	Normal	Scheduled	pod/nginx-8cdb45866-9p8ws	Successfully assigned default/nginx-8cdb45866-9p8ws to ubuntu
63s	Normal	Pulling	pod/nginx-8cdb45866-9p8ws	Pulling image "nginx:latest"
57s	Normal	Pulled	pod/nginx-8cdb45866-9p8ws	Successfully pulled image "nginx:latest"
57s	Normal	Created	pod/nginx-8cdb45866-9p8ws	Created container nginx
57s	Normal	Started	pod/nginx-8cdb45866-9p8ws	Started container nginx

```

5s          Warning   Unhealthy           pod/nginx-8cdb45866-9p8ws    Liveness probe
failed: Get http://10.32.0.6:80/: dial tcp 10.32.0.6:80: connect: connection refused

user@ubuntu:~/hc$

```

What happened?

Events reported by the event stream are not as granular as those provided by the describe, try it:

```

user@ubuntu:~/hc$ kubectl describe pod \
$(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}') | grep -A 15 Events

Events:
  Type            Reason          Age           From              Message
  ----            -
  Normal          Scheduled        <unknown>     default-scheduler Successfully assigned default/nginx-
8cdb45866-9p8ws to ubuntu
  Warning         Unhealthy        19s (x3 over 39s) kubelet, ubuntu    Liveness probe failed: Get
http://10.32.0.6:80/: dial tcp 10.32.0.6:80: connect: connection refused
  Normal          Killing          19s          kubelet, ubuntu    Container nginx failed liveness
probe, will be restarted
  Normal          Pulling          18s (x2 over 87s) kubelet, ubuntu    Pulling image "nginx:latest"
  Normal          Pulled           17s (x2 over 81s) kubelet, ubuntu    Successfully pulled image
"nginx:latest"
  Normal          Created          17s (x2 over 81s) kubelet, ubuntu    Created container nginx
  Normal          Started          17s (x2 over 81s) kubelet, ubuntu    Started container nginx

user@ubuntu:~/hc$

```

As you can see the Liveness probe failed at one point, but recovered. The nginx container in the pod was created, started, found unhealthy, killed, created and started again.

Remove the related resources:

```

user@ubuntu:~/jobs$ kubectl delete deploy/nginx

deployment.apps "nginx" deleted

user@ubuntu:~/jobs$

```

9. Creating a Job

In a previous lab we saw that running a pod standalone works but without an RS the pod will not restart if it crashes. Unfortunately, if we run a batch job in a pod with an RS and the pod completes the task, the RS will start the pod again.

What if we want a pod that runs only once, however, if it or the node it is running on fails before the pod completes successfully, we want the pod to be started again until it does complete successfully. Kubernetes provides a **Job** type for this scenario.

A Job is like an RC/RS that ensures that a pod runs once to completion. Imagine we want to calculate Pi. Not twice, not half of a time, but precisely once. A job would be the perfect way to run a container that calculates Pi. Enter this sample job config to compute Pi:

```

user@ubuntu:~/hc$ cd ~

user@ubuntu:~$ mkdir ~/jobs && cd ~/jobs/

user@ubuntu:~/jobs$ nano myjob.yaml && cat myjob.yaml

apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    metadata:

```

```

name: pi
spec:
  containers:
  - name: pi
    image: perl
    command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
    restartPolicy: Never

```

```
user@ubuntu:~/jobs$
```

The config uses apiVersion "batch/v1". The kind of object we will create is a Job. The Job will have the name "pi", as per the metadata. The template for the pod the Job we'll create must have a name pi.

The spec for the pod uses a single perl container which will run the command that computes pi. We also set the restart policy to Never.

Now try running your Job:

```
user@ubuntu:~/jobs$ kubectl apply -f myjob.yaml
```

```
job.batch/pi created
```

```
user@ubuntu:~/jobs$
```

Examine the job:

```
user@ubuntu:~/jobs$ kubectl describe job/pi
```

```

Name:          pi
Namespace:     default
Selector:      controller-uid=bc8c99dc-2484-4879-aeae-0b00bccaa29c
Labels:        controller-uid=bc8c99dc-2484-4879-aeae-0b00bccaa29c
               job-name=pi
Annotations:   kubectl.kubernetes.io/last-applied-configuration:
               {"apiVersion":"batch/v1","kind":"Job","metadata":{"annotations":
               {},"name":"pi","namespace":"default"},"spec":{"template":{"metadata":{"nam...
Parallelism:   1
Completions:   1
Start Time:    Wed, 08 Jan 2020 13:48:18 -0800
Pods Statuses: 1 Running / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  controller-uid=bc8c99dc-2484-4879-aeae-0b00bccaa29c
          job-name=pi
  Containers:
    pi:
      Image:          perl
      Port:           <none>
      Host Port:      <none>
      Command:
        perl
        -Mbignum=bpi
        -wle
        print bpi(2000)
      Environment:    <none>
      Mounts:          <none>
      Volumes:         <none>
Events:
  Type     Reason             Age   From             Message
  ----     -
  Normal   SuccessfulCreate   7s    job-controller   Created pod: pi-2l6j6

```

The **kubectl apply** subcommand processes the job request and runs our pod. Displaying the Job description shows us the name of the pod that ran the Job. We can now dump the logs for the pod to see the result:

```
user@ubuntu:~/jobs$ kubectl logs $(kubectl get jobs -o name)
```

```
3.1415926535897932384626433832795028841971693993751058209749445923078164062862089986280348253421
170679821480865132823066470938446095505822317253594081284811174502841027019385211055596446229489
549303819644288109756659334461284756482337867831652712019091456485669234603486104543266482133936
072602491412737245870066063155881748815209209628292540917153643678925903600113305305488204665213
841469519415116094330572703657595919530921861173819326117931051185480744623799627495673518857527
248912279381830119491298336733624406566430860213949463952247371907021798609437027705392171762931
767523846748184676694051320005681271452635608277857713427577896091736371787214684409012249534301
465495853710507922796892589235420199561121290219608640344181598136297747713099605187072113499999
983729780499510597317328160963185950244594553469083026425223082533446850352619311881710100031378
387528865875332083814206171776691473035982534904287554687311595628638823537875937519577818577805
321712268066130019278766111959092164201989380952572010654858632788659361533818279682303019520353
018529689957736225994138912497217752834791315155748572424541506959508295331168617278558890750983
817546374649393192550604009277016711390098488240128583616035637076601047101819429555961989467678
374494482553797747268471040475346462080466842590694912933136770289891521047521620569660240580381
501935112533824300355876402474964732639141992726042699227967823547816360093417216412199245863150
302861829745557067498385054945885869269956909272107975093029553211653449872027559602364806654991
198818347977535663698074265425278625518184175746728909777727938000816470600161452491921732172147
723501414419735685481613611573525521334757418494684385233239073941433345477624168625189835694855
620992192221842725502542568876717904946016534668049886272327917860857843838279679766814541009538
837863609506800642251252051173929848960841284886269456042419652850222106611863067442786220391949
450471237137869609563643719172874677646575739624138908658326459958133904780275901
```

```
user@ubuntu:~/jobs$
```

By default, a Job is complete when one Pod runs to successful completion. You can also specify that this needs to happen multiple times by specifying Job spec key *"completions"* with a value greater than 1. You can suggest how many pods should run concurrently by setting Job spec key *"parallelism"* to the number of pods you would like to have running concurrently (the value defaults to *"completions"*.) The parallelism key is just a hint and the Job may run fewer or more concurrent pods.

Jobs are complementary to Deployments. A Deployment manages pods which are not expected to terminate (e.g. web servers,) and a Job manages pods that are expected to terminate (e.g. batch jobs.)

Complete job pods and the jobs themselves remain after completion, so their logs can be called at any time.

Check the pods and jobs on your cluster:

```
user@ubuntu:~/jobs$ kubectl get pods,jobs
```

NAME	READY	STATUS	RESTARTS	AGE
pod/pi-2l6j6	0/1	Completed	0	74s

NAME	COMPLETIONS	DURATION	AGE
job.batch/pi	1/1	52s	74s

```
user@ubuntu:~/jobs$
```

It is up to you to remove the job at your discretion. The only way to remove the job pod is by removing the job itself, so When you are finished exploring remove the Job:

```
user@ubuntu:~/jobs$ kubectl delete job pi
```

```
job.batch "pi" deleted
```

```
user@ubuntu:~/jobs$
```

```
user@ubuntu:~/jobs$ kubectl get deploy,rs,pods,job
```

```
No resources found in default namespace.
```

```
user@ubuntu:~/jobs$ cd ~
```

```
user@ubuntu:~$
```

Congratulations you have completed the lab!

Copyright (c) 2013-2020 RX-M LLC, Cloud Native Consulting, all rights reserved