

Kubernetes

Lab 8 – Metrics

Starting from Kubernetes 1.8, resource usage metrics, such as container CPU and memory usage, are available in Kubernetes through the Metrics API. These metrics can be either accessed directly by a user with `kubectl top`, or used by a controller in the cluster, such as the Horizontal Pod Autoscaler.

The Metrics API reports the current resource utilization of a node or pod or set thereof. The metrics API is designed to expose the current resource utilization only. Those wishing to view a history of resource utilization should use a monitoring tool and database such as Prometheus or InfluxDB/Grafana.

Kubernetes 1.7 introduced the aggregator API feature. This allows add on services to extend the base K8s API. The Metrics API uses the aggregator feature and exposes an API available at `/apis/metrics.k8s.io/`, offering the same security, scalability and reliability guarantees as the rest of the api-server hosted end points.

The Metrics server was first released in September 2017 and is not installed automatically by all K8s installers. The Metrics Server does not run as a module within the Controller Manager, rather it runs as a stand alone deployment, typically in the `kube-system` namespace. The Metrics Server requires an SA with appropriate RBAC roles so that it can communicate with the api-server and extend the K8s API.

The Metrics Server must also be equipped with certificates/keys that will allow it to connect to each of the Kubelets in the system. Kubelets expose a `/metrics` end point, the Metrics Server scrapes this endpoint for metrics data regularly, storing the results in memory. When in bound requests for metrics arrive at the api-server the metrics server is called to answer them.

In this lab we'll install the Metrics Server and explore the metrics enabled features of Kubernetes. Delete all user defined deployments, services, and other resources before starting.

1. Explore the Metrics Server repository

Basic Metrics Server deployment manifests are available in the `/deploy` directory in the Metrics Server source repo:

<https://github.com/kubernetes-incubator/metrics-server>

The Metrics Server repository includes a set of manifests that create all of the cluster resources needed to run the Metrics Server. Clone the Metrics Server repository from GitHub:

```
user@ubuntu:~$ git clone --depth 1 https://github.com/kubernetes-incubator/metrics-server

Cloning into 'metrics-server'...
remote: Enumerating objects: 88, done.
remote: Counting objects: 100% (88/88), done.
remote: Compressing objects: 100% (74/74), done.
remote: Total 88 (delta 11), reused 39 (delta 5), pack-reused 0
Unpacking objects: 100% (88/88), done.
Checking connectivity... done.

user@ubuntu:~$
```

Now change into the repository directory and display the manifests in the `deploy/1.8+` directory:

```
user@ubuntu:~$ cd metrics-server/

user@ubuntu:~/metrics-server$ ls -l

total 152
drwxrwxr-x  3 user user  4096 Jan 24 15:44 cmd
-rw-rw-r--  1 user user   148 Jan 24 15:44 code-of-conduct.md
-rw-rw-r--  1 user user   418 Jan 24 15:44 CONTRIBUTING.md
drwxrwxr-x  6 user user  4096 Jan 24 15:44 deploy
```

```

-rw-rw-r-- 1 user user 3364 Jan 24 15:44 go.mod
-rw-rw-r-- 1 user user 84975 Jan 24 15:44 go.sum
drwxrwxr-x 2 user user 4096 Jan 24 15:44 hack
-rw-rw-r-- 1 user user 11357 Jan 24 15:44 LICENSE
-rw-rw-r-- 1 user user 4020 Jan 24 15:44 Makefile
-rw-rw-r-- 1 user user 201 Jan 24 15:44 OWNERS
-rw-rw-r-- 1 user user 177 Jan 24 15:44 OWNERS_ALIASES
drwxrwxr-x 11 user user 4096 Jan 24 15:44 pkg
-rw-rw-r-- 1 user user 5305 Jan 24 15:44 README.md
-rw-rw-r-- 1 user user 541 Jan 24 15:44 SECURITY_CONTACTS
drwxrwxr-x 2 user user 4096 Jan 24 15:44 test

```

```
user@ubuntu:~/metrics-server$ ls -l deploy/
```

```

total 16
drwxrwxr-x 2 user user 4096 Jan 24 15:44 1.7
drwxrwxr-x 2 user user 4096 Jan 24 15:44 1.8+
drwxrwxr-x 2 user user 4096 Jan 24 15:44 docker
drwxrwxr-x 2 user user 4096 Jan 24 15:44 minikube

```

```
user@ubuntu:~/metrics-server$ ls -l deploy/1.8+/
```

```

total 28
-rw-rw-r-- 1 user user 397 Jan 24 15:44 aggregated-metrics-reader.yaml
-rw-rw-r-- 1 user user 303 Jan 24 15:44 auth-delegator.yaml
-rw-rw-r-- 1 user user 324 Jan 24 15:44 auth-reader.yaml
-rw-rw-r-- 1 user user 298 Jan 24 15:44 metrics-apiservice.yaml
-rw-rw-r-- 1 user user 1183 Jan 24 15:44 metrics-server-deployment.yaml
-rw-rw-r-- 1 user user 297 Jan 24 15:44 metrics-server-service.yaml
-rw-rw-r-- 1 user user 532 Jan 24 15:44 resource-reader.yaml

```

```
user@ubuntu:~/metrics-server$
```

The Metrics Server deployment uses several manifests. Generate a listing of the types of resources that will be created:

```
user@ubuntu:~/metrics-server$ kubectl apply -f deploy/1.8+/- --dry-run
```

```

clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created (dry run)
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created (dry run)
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created (dry run)
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created (dry run)
serviceaccount/metrics-server created (dry run)
deployment.apps/metrics-server created (dry run)
service/metrics-server created (dry run)
clusterrole.rbac.authorization.k8s.io/system:metrics-server created (dry run)
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created (dry run)

```

```
user@ubuntu:~/metrics-server$
```

Though a range of resources are defined, you can see that most are RBAC elements:

- deploy/1.8+/aggregated-metrics-reader.yaml:kind: ClusterRole
- deploy/1.8+/auth-delegator.yaml:kind: ClusterRoleBinding
- deploy/1.8+/auth-reader.yaml:kind: RoleBinding
- deploy/1.8+/metrics-server-deployment.yaml:kind: ServiceAccount
- deploy/1.8+/resource-reader.yaml:kind: ClusterRole
- deploy/1.8+/resource-reader.yaml:kind: ClusterRoleBinding

The manifests create two ClusterRoles. Display the rules defined:

```
user@ubuntu:~/metrics-server$ grep rules -A21 deploy/1.8+/*
```

```

deploy/1.8+/aggregated-metrics-reader.yaml:rules:
deploy/1.8+/aggregated-metrics-reader.yaml-- apiGroups: ["metrics.k8s.io"]
deploy/1.8+/aggregated-metrics-reader.yaml- resources: ["pods", "nodes"]

```

```

deploy/1.8+/aggregated-metrics-reader.yaml- verbs: ["get", "list", "watch"]
--
deploy/1.8+/resource-reader.yaml:rules:
deploy/1.8+/resource-reader.yaml-- apiGroups:
deploy/1.8+/resource-reader.yaml- - ""
deploy/1.8+/resource-reader.yaml- resources:
deploy/1.8+/resource-reader.yaml- - pods
deploy/1.8+/resource-reader.yaml- - nodes
deploy/1.8+/resource-reader.yaml- - nodes/stats
deploy/1.8+/resource-reader.yaml- - namespaces
deploy/1.8+/resource-reader.yaml- - configmaps
deploy/1.8+/resource-reader.yaml- verbs:
deploy/1.8+/resource-reader.yaml- - get
deploy/1.8+/resource-reader.yaml- - list
deploy/1.8+/resource-reader.yaml- - watch
deploy/1.8+/resource-reader.yaml----
deploy/1.8+/resource-reader.yaml-apiVersion: rbac.authorization.k8s.io/v1
deploy/1.8+/resource-reader.yaml-kind: ClusterRoleBinding
deploy/1.8+/resource-reader.yaml-metadata:
deploy/1.8+/resource-reader.yaml- name: system:metrics-server
deploy/1.8+/resource-reader.yaml-roleRef:
deploy/1.8+/resource-reader.yaml- apiGroup: rbac.authorization.k8s.io
deploy/1.8+/resource-reader.yaml- kind: ClusterRole
deploy/1.8+/resource-reader.yaml- name: system:metrics-server

user@ubuntu:~/metrics-server$

```

The aggregated-metrics-reader provides access to the pod information within the metrics.k8s.io API group. The resource-reader provides access to information methods for pods, nodes, nodes/stats and namespaces across all API groups and information access to deployments within the extensions API group.

The ClusterRoleBinding objects bind the ClusterRoles to the metrics-server ServiceAccount created in the kube-system namespace. The one RoleBinding binds the preexisting extension-apiserver-authentication-reader Role to the Metric Server's service account:

```

user@ubuntu:~/metrics-server$ cat deploy/1.8+/auth-reader.yaml

```

```

---
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: metrics-server-auth-reader
  namespace: kube-system
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: extension-apiserver-authentication-reader
subjects:
- kind: ServiceAccount
  name: metrics-server
  namespace: kube-system

```

```

user@ubuntu:~/metrics-server$ kubectl get role -n kube-system extension-apiserver-
authentication-reader

NAME                                     AGE
extension-apiserver-authentication-reader 4h25m

user@ubuntu:~/metrics-server$

```

The extension-apiserver-authentication-reader allows the addon API server to read the client CA file from the extension-apiserver-authentication ConfigMap in the kube-system namespace. This is required in order for the Metrics addon API to delegate authentication and authorization requests to the main Kubernetes API server.

The remaining three manifests support the Metrics Server operation:

- metrics-apiservice.yaml:kind: APIService

- metrics-server-deployment.yaml:kind: Deployment
- metrics-server-service.yaml:kind: Service

The APIService manifest defines the metric-server service API extension to the core Kubernetes API:

```
user@ubuntu:~/metrics-server$ cat deploy/1.8+/metrics-apiservice.yaml
```

```
---
apiVersion: apiregistration.k8s.io/v1beta1
kind: APIService
metadata:
  name: v1beta1.metrics.k8s.io
spec:
  service:
    name: metrics-server
    namespace: kube-system
  group: metrics.k8s.io
  version: v1beta1
  insecureSkipTLSVerify: true
  groupPriorityMinimum: 100
  versionPriority: 100
```

```
user@ubuntu:~/metrics-server$
```

Ad hoc API extensions make use of the Kubernetes API aggregation function. Kubernetes API aggregation enables the dynamic installation of additional Kubernetes-style APIs in a cluster. These can either be pre-built, existing 3rd party solutions, such as service-catalog, or user-created APIs like apiserver-builder.

The aggregation layer runs in-process with the kube-apiserver. Until an extension resource is registered, the aggregation layer will do nothing. To register an API an APIService object is used to "claim" a URL path in the Kubernetes API. At this point, the aggregation layer will proxy anything sent to that API path to the registered APIService.

In this case the metrics-server claims the `/apis/metrics.k8s.io/` path.

The metrics server deployment is fairly simple and just runs the `k8s.gcr.io/metrics-server-amd64:v0.3.6` image. The Metrics Server service is also simple, creating a "metrics-server" service that accepts traffic on port 443.

Examine any of the yaml manifests you find interesting.

2. Deploy the Metrics Server

N.B. If you have a multi-node cluster, make sure you deploy the metrics server to your master node. Use `kubectl cordon <worker node>` and remove the master taint if you have it with `kubectl taint nodes $(hostname) node-role.kubernetes.io/master-` to ensure this lab's pod run on your master node.

We can launch the Metrics Server by creating all of the resources in the `deploy/1.8+` directory by providing `kubectl apply -f` with the directory path. Try it:

```
user@ubuntu:~/metrics-server$ kubectl apply -f deploy/1.8+/

clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
serviceaccount/metrics-server created
deployment.apps/metrics-server created
service/metrics-server created
clusterrole.rbac.authorization.k8s.io/system:metrics-server created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created

user@ubuntu:~/metrics-server$
```

```
user@ubuntu:~/metrics-server$ kubectl get all -n kube-system -l k8s-app=metrics-server
```

NAME	READY	STATUS	RESTARTS	AGE
pod/metrics-server-694db48df9-lzm99	1/1	Running	0	43s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/metrics-server	1/1	1	1	43s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/metrics-server-694db48df9	1	1	1	43s

```
user@ubuntu:~/metrics-server$
```

Great, the Metrics Server is deployed and running! Now try a metrics api request:

```
user@ubuntu:~/metrics-server$ curl -k https://localhost:6443/apis/metrics.k8s.io/
```

```
{
  "kind": "Status",
  "apiVersion": "v1",
  "metadata": {
  },
  "status": "Failure",
  "message": "forbidden: User \"system:anonymous\" cannot get path \"/apis/metrics.k8s.io/\"",
  "reason": "Forbidden",
  "details": {
  },
  "code": 403
}
```

```
user@ubuntu:~/metrics-server$
```

We cannot connect to the apiserver without a user identity, let's use an existing one:

```
user@ubuntu:~/metrics-server$ kubectl get sa
```

NAME	SECRETS	AGE
default	1	24h

```
user@ubuntu:~/metrics-server$
```

Get the token and certificate from the ServiceAccount's token secret for use in your API requests.

This will require `jq` for JSON parsing on the command line, install it:

```
user@ubuntu:~/metrics-server$ sudo apt-get install jq -y
...
user@ubuntu:~/metrics-server$
```

Start by setting the `SERVICE_ACCOUNT` variable:

```
user@ubuntu:~/metrics-server$ SERVICE_ACCOUNT=default
user@ubuntu:~/metrics-server$
```

```
user@ubuntu:~/metrics-server$ SECRET=$(kubectl get serviceaccount ${SERVICE_ACCOUNT} -o json \
| jq -Mr '.secrets[] | select(contains("token"))') && echo $SECRET

default-token-ksft2

user@ubuntu:~/metrics-server$
```

[illegible]

```
user@ubuntu:~/metrics-server$ kubectl get secret ${SECRET} -o json \
| jq -Mr '.data["ca.crt"]' | base64 -d > /tmp/ca.crt

user@ubuntu:~/metrics-server$
```

```
user@ubuntu:~/metrics-server$ curl -k https://localhost:6443/apis/metrics.k8s.io/ \
--header "Authorization: Bearer $TOKEN" --cacert /tmp/ca.crt
```

```
{
  "kind": "APIGroup",
  "apiVersion": "v1",
  "name": "metrics.k8s.io",
  "versions": [
    {
      "groupVersion": "metrics.k8s.io/v1beta1",
      "version": "v1beta1"
    }
  ],
  "preferredVersion": {
    "groupVersion": "metrics.k8s.io/v1beta1",
    "version": "v1beta1"
  }
}
```

```
user@ubuntu:~/metrics-server$
```

Page 6/21 © Copyright Thursday, Jan 30, 2020, 8:58 AM RX-M LLC

3. Test the Metrics Server

The `kubectl top` command can be used to display metrics for pods or nodes. Try it:

```
user@ubuntu:~/metrics-server$ kubectl top node

error: metrics not available yet

user@ubuntu:~/metrics-server$
```

Hmm, no metrics. Try checking the metrics server log to see if there are any issues indicated by the metrics server pod:

```
user@ubuntu:~/metrics-server$ kubectl get pods -n kube-system -l k8s-app=metrics-server

NAME                                READY   STATUS    RESTARTS   AGE
metrics-server-694db48df9-lzm99    1/1     Running   0           78s

user@ubuntu:~/metrics-server$ kubectl -n kube-system logs metrics-server-694db48df9-lzm99

I0124 23:49:33.028127       1 serving.go:312] Generated self-signed cert (/tmp/apiserver.crt, /tmp/apiserver.key)
I0124 23:49:33.303878       1 secure_serving.go:116] Serving securely on [::]:4443
E0124 23:50:45.557373       1 manager.go:111] unable to fully collect metrics: unable to fully scrape metrics from source kubelet_summary:ubuntu: unable to fetch metrics from Kubelet ubuntu (ubuntu): Get https://ubuntu:10250/stats/summary?only_cpu_and_memory=true: dial tcp 66.96.162.149:10250: connect: connection refused
E0124 23:51:19.105292       1 reststorage.go:135] unable to fetch node metrics for node "ubuntu": no metrics known for node
E0124 23:51:45.520002       1 manager.go:111] unable to fully collect metrics: unable to fully scrape metrics from source kubelet_summary:ubuntu: unable to fetch metrics from Kubelet ubuntu (ubuntu): Get https://ubuntu:10250/stats/summary?only_cpu_and_memory=true: dial tcp 66.96.162.149:10250: connect: connection refused

user@ubuntu:~/metrics-server$
```

One issue here is that the metrics server cannot communicate with any of the kubelets for this host. In the example above, the hostname "ubuntu" is routing to an IP address that is rejecting the connection, meaning the DNS configuration in this environment is not configured correctly.

Even if it could reach the host, another problem is that the Metrics server requests its metrics from the Kubelets but rejects the kubelet responses. This is because the Metrics Server attempts to use the api-server CA cert to verify the Kubelet's response. In the Kubeadm case today the Kubelet uses an api-server CA signed cert to reach out to the api-server but when listening on its own API port (where the /metrics endpoint we want is located) it uses a self signed certificate. This causes the Metrics Server to reject the TLS session upon connection. Existing workarounds will suggest using `--kubelet-insecure-tls` to tell the Metrics Server to not check the Kubelet's TLS cert. This approach is not ideal for production deployments.

3a. Force the kubelet to request a new certificate from the cluster

To fix this you will need to regenerate the kubelet's certificate using the cluster CA certificate. This is possible by forcing the node kubelet to go through the TLS bootstrap process again to request a new serving certificate. The serving certificate bears the cluster CA's signature and also includes the hostname and IP addresses of the kubelet as subject alternative names. Any services that need to verify the kubelet's identities will be able to do so using a new serving certificate.

To begin, retrieve the bootstrap token from the cluster. This token is what your node's kubelet used initially to communicate with the API server. This bootstrap token is assigned to the system:bootstrappers group, which gives it all the permissions it needs to connect to the API server and request a certificate. These tokens are backed by secrets in the cluster.

Retrieve the bootstrap token from the cluster's secrets:

```
user@ubuntu:~/metrics-server$ kubectl get secrets -n kube-system | grep bootstrap

bootstrap-signer-token-mnmtt                                kubernet.es.io/service-account-token    3
24h
bootstrap-token-pllvту                                       bootstrap.kubernet.es.io/token          7
```

24h

```
user@ubuntu:~/metrics-server$ export BOOTSTRAPSECRET=bootstrap-token-pllvtu

user@ubuntu:~/metrics-server$ kubectl get secret -n kube-system $BOOTSTRAPSECRET -o yaml

apiVersion: v1
data:
  auth-extra-groups: c3lzdGVtOmJvb3RzdHJhcHB1cnM6a3ViZWFKbTpkZWZhdWx0LW5vZGUtdG9rZW4=
  expiration: MjAyMC0wMS0yNVQxNjoxMDoyMi0wODowMA==
  token-id: YzJpaGo1
  token-secret: bDZhOG5qOHJkeHBsZXpzOQ==
  usage-bootstrap-authentication: dHJ1ZQ==
  usage-bootstrap-signing: dHJ1ZQ==
kind: Secret
metadata:
  creationTimestamp: "2020-01-25T00:10:22Z"
  name: bootstrap-token-c2ihj5
  namespace: kube-system
  resourceVersion: "26242"
  selfLink: /api/v1/namespaces/kube-system/secrets/bootstrap-token-c2ihj5
  uid: 082db190-155c-4a3f-8d9d-e62a2080427d
type: bootstrap.kubernetes.io/token

user@ubuntu:~/metrics-server$
```

The token-id and token-secret are needed to present a complete bootstrap token.

Using `-o jsonpath` to extract specific values from the secret, retrieve the `token-id` and the `token-secret`. To output the values in plain text, pipe the output into `base64 --decode`:

```
user@ubuntu:~/metrics-server$ kubectl get secret -n kube-system $BOOTSTRAPSECRET -o
jsonpath='{.data.token-id}' \
| base64 --decode && echo

c2ihj5

user@ubuntu:~/metrics-server$ kubectl get secret -n kube-system $BOOTSTRAPSECRET -o
jsonpath='{.data.token-secret}' \
| base64 --decode && echo

16a8nj8rdxplezs9

user@ubuntu:~/metrics-server$
```

The proper format for a token is `token-id.token-secret`. When properly formatted, the complete bootstrap token for this cluster is: `c2ihj5.16a8nj8rdxplezs9`

N.B. In Kubeadm clusters like ours, bootstrap tokens are only valid for 24 hours. If your bootstrap token is over 24 hours old, you can use kubeadm to create a new token with `kubeadm token create`. This command outputs a complete bootstrap token to the terminal and creates a corresponding secret in the cluster.

```
user@ubuntu:~$ kubeadm token create

ezgumy.13p5h3hn4xsd9u2k

user@ubuntu:~$
```

If this is how you retrieved the bootstrap token, make sure to use whatever output is put in the terminal for the following steps.

Now you need to create a bootstrap kubelet configuration. This configuration tells the Kubelet to re-register itself with the cluster. First you need to see the cluster's connection information by reading the current kubelet configuration at `/etc/kubernetes/kubelet.conf`:


```

user@ubuntu:~/metrics-server$ sudo head /etc/kubernetes/kubelet.conf

apiVersion: v1
clusters:
- cluster:
    certificate-authority-data:
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSB0tLS0tCk1JSUN5RENDQWJDZ0F3SUJBZ0lCQURBTk1na3Foa2lHOXcwQkFRc0ZBREWFw
TVJNd0VRWURWUWFERXdwcmRXSmwKY201bGRHVnpNQjRYFRjZ01ERXlNek16TkRveU5Wb1hEVE13TURFeU1ESXp0RFV5T1Zv
d0ZURVRnQkVHQTFRVQpBeE1LYTNwaVpYSnVwVjJ5Y3pDQ0FTSXdEUVlKS29aSWh2Y05BUUVCQlFBRGdnRVBBRENDQVFvQ2dn
RUJBSjdJc2NUGpQeUxTZWZldFQvZE9KYlllV1RKMWE0eEovSngzbFByV3ovcDFUcEVsUmZVemFwVFR5bUw2Tk1MFd2VFAK
M1I3ZXdsakc0WwXZWXRKN0tKRk1JR2N4UFBrzdHV3AxSUR1YnRrRjhtYTVXU2NNNHple1I0K3FEbmYyS2ZUNWpVNldscWt0
WEUzcDR2YXhBTERXNHua0NPS1pZQXE1RlFWb1owQTR1VzU5bkY1ZFJ5RGNXMkw3dWRQYVJvQ3A2CmZlbnR5PSUo4TUtsWjQ3
UHJkWFhCWGU1U0lYZURVZGtJUKwrRXdRbnpawjUUL05MSkYwbXpEeTVrYzd3SHIrMTEKaStCWDdPaDNUYzFtUDlVaVRha2V1
cTA3M2JIN2c0TGFuR0ExRXkvK1VWZVIyTGg5V3grOUJJQjBSaFJQZ2FNdAoZS1doL0I1ME9DU3FFVVMzZ204Q0F3RUFBYU1q
TUNFd0RnWURWUjBQOVFILOjBUURBZ0trTUE4R0ExVWRFd0VCCi93UUZUNQU1CQWY4d0RRWUplb1pJaHZjTkFRRUxUUUFEZ2dF
QkFIQVJPNKxwMEJ5NmpzTS9MWDZaThITnJKTU8KdVl0ejk5aTc0OXlHdEJOQjJHd1FJMHd5NzJlNGFvMWU2MG5Xe1h1YXcw
QTJBWldNbTIwVzZjM1t1c1JpV1dCNAp2WDB2UTU0cGE5YzE5dmluUTk4NXVwNmNmNnFyYVRZQzNjWWtuTDVxSnZsN1RwdUpV
QkQyUkhIVDFsS0lxSG5XCnpDVXhXZjdpVk55dmxwMFZmTlhfUUhYK0phN05jSTARA0R0ZHFqVFQxZ21MdmM4b0dUb1FtN0tu
M0d2My9aRlMKem1yK0JSQjNIYVg1ZUJoVHVkRFP6WDNQYTZtc28yYmtSVUR4Y2FWcHB1bmVmdWZlbnR5PSUo4TUtsWjQ3
Ywo3Q1l0bGVVMXVFeTF0U1BrdjFhNHpYc1pXTDBY0XpVdGJ4b0JONUZ1V1B1K1Z0RlNPd0V3Q25BVDVSZ0KLS0tLS1FTkQg
Q0VSVElGSUNBVEUtLS0tLQo=
    server: https://192.168.229.132:6443
    name: default-cluster
contexts:
- context:
    cluster: default-cluster
    namespace: default

user@ubuntu:~/metrics-server$

```

We will use the value of the `server` key under `-cluster:` to inform the new bootstrap config how to connect to the cluster.

Now we can construct the bootstrap configuration using `kubect1`. First, add the cluster connection and certificate authority information to the file using `kubect1 config set-cluster`. Be sure to substitute the `--server` flag with the value of `server:` from the kubelet.conf file:

```

user@ubuntu:~/metrics-server$ kubect1 config set-cluster bootstrap \
--kubeconfig=bootstrap-kubelet.conf \
--certificate-authority='/etc/kubernetes/pki/ca.crt' \
--server='https://192.168.229.132:6443'

```

Cluster "bootstrap" set.

```
user@ubuntu:~/metrics-server$
```

Next, associate the bootstrap token you retrieved earlier to a user called `kubelet-bootstrap`. Be sure to substitute the value of `--token` with the token you retrieved or created earlier:

```

user@ubuntu:~/metrics-server$ kubect1 config set-credentials kubelet-bootstrap \
--kubeconfig=bootstrap-kubelet.conf \
--token=c2ihj5.l6a8nj8rdxplezs9

```

User "kubelet-bootstrap" set.

```
user@ubuntu:~/metrics-server$
```

Now, set a context that associates the `kubelet-bootstrap` user to the cluster connection information:

```

ubuntu@ubuntu:~$ kubect1 config set-context bootstrap \
--kubeconfig=bootstrap-kubelet.conf \
--user=kubelet-bootstrap --cluster=bootstrap

```

Context "bootstrap" created.

```
user@ubuntu:~/metrics-server$
```

Finally, create the file by switching to the bootstrap context with `kubectl config use-context` :

```
user@ubuntu:~/metrics-server$ kubectl config use-context bootstrap \
--kubeconfig=bootstrap-kubelet.conf

Switched to context "bootstrap".

user@ubuntu:~/metrics-server$
```

The file is now available. Use `cat` to view the contents.

```
user@ubuntu:~/metrics-server$ cat bootstrap-kubelet.conf

apiVersion: v1
clusters:
- cluster:
    certificate-authority: /etc/kubernetes/pki/ca.crt
    server: https://192.168.229.132:6443
    name: bootstrap
contexts:
- context:
    cluster: bootstrap
    user: kubelet-bootstrap
    name: bootstrap
current-context: bootstrap
kind: Config
preferences: {}
users:
- name: kubelet-bootstrap
  user:
    token: c2ihj5.16a8nj8rdxplezs9

user@ubuntu:~/metrics-server$
```

Copy `bootstrap-kubelet.conf` to `/etc/kubernetes/` :

```
user@ubuntu:~/metrics-server$ sudo cp bootstrap-kubelet.conf /etc/kubernetes/bootstrap-
kubelet.conf

user@ubuntu:~/metrics-server$ sudo cat /etc/kubernetes/bootstrap-kubelet.conf

apiVersion: v1
clusters:
- cluster:
    certificate-authority: /etc/kubernetes/pki/ca.crt
    server: https://192.168.229.132:6443
    name: bootstrap
contexts:
- context:
    cluster: bootstrap
    user: kubelet-bootstrap
    name: bootstrap
current-context: bootstrap
kind: Config
preferences: {}
users:
- name: kubelet-bootstrap
  user:
    token: c2ihj5.16a8nj8rdxplezs9

user@ubuntu:~/metrics-server$
```

To ensure the kubelet requests a new signed serving certificate, we need to adjust the kubelet arguments to use server certificate rotation and force the cluster to distribute a new certificate to the kubelet. Add `--rotate-server-certificates` to your kubelet. For kubeadm-bootstrapped clusters, you can add it to the `kubeadm-flags.env` file:

```
user@ubuntu:~/metrics-server$ sudo nano /var/lib/kubelet/kubeadm-flags.env

user@ubuntu:~/metrics-server$ sudo cat /var/lib/kubelet/kubeadm-flags.env

KUBELET_KUBEADM_ARGS="--cgroup-driver=cgroupfs --network-plugin=cni --pod-infra-container-
image=k8s.gcr.io/pause:3.1 --rotate-server-certificates"

user@ubuntu:~/metrics-server$
```

N.B. If your cluster does not have a `kubeadm-flags.env` file, then edit the kubelet systemd service file directly and place `--rotate-server-certificates` inline with the `ExecStart` line pointing to the kubelet binary.

Next, we need to force the kubelet to use the bootstrap configuration by removing the existing `kubelet.conf` file. Renaming the old file will suffice:

```
user@ubuntu:~/metrics-server$ sudo mv /etc/kubernetes/kubelet.conf /etc/kubernetes/kubelet.old
```

Now restart the kubelet:

```
user@ubuntu:~/metrics-server$ sudo systemctl restart kubelet

user@ubuntu:~/metrics-server$ sudo systemctl status kubelet

● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/lib/systemd/system/kubelet.service; enabled; vendor preset: enabled)
   Drop-In: /etc/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Fri 2020-01-24 16:19:01 PST; 2s ago
     Docs: https://kubernetes.io/docs/home/
   Main PID: 125944 (kubelet)
    Tasks: 15
   Memory: 34.5M
      CPU: 420ms
   CGroup: /system.slice/kubelet.service
            └─125944 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-
kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --config=/var/lib/kubelet/config.yaml --
cgroup-driver=cgroupfs

Jan 24 16:19:02 ubuntu kubelet[125944]: I0124 16:19:02.434257 125944 reconciler.go:207]
operationExecutor.VerifyControllerAttachedVolume started for volume "etcd-certs" (UniqueName:
"kubernetes.io/host-
Jan 24 16:19:02 ubuntu kubelet[125944]: I0124 16:19:02.434268 125944 reconciler.go:207]
operationExecutor.VerifyControllerAttachedVolume started for volume "weavedb" (UniqueName:
"kubernetes.io/host-pat
Jan 24 16:19:02 ubuntu kubelet[125944]: I0124 16:19:02.434282 125944 reconciler.go:207]
operationExecutor.VerifyControllerAttachedVolume started for volume "dbus" (UniqueName:
"kubernetes.io/host-path/d
Jan 24 16:19:02 ubuntu kubelet[125944]: I0124 16:19:02.434295 125944 reconciler.go:207]
operationExecutor.VerifyControllerAttachedVolume started for volume "config-volume" (UniqueName:
"kubernetes.io/co
Jan 24 16:19:02 ubuntu kubelet[125944]: I0124 16:19:02.434308 125944 reconciler.go:207]
operationExecutor.VerifyControllerAttachedVolume started for volume "config-volume" (UniqueName:
"kubernetes.io/co
Jan 24 16:19:02 ubuntu kubelet[125944]: I0124 16:19:02.434313 125944 reconciler.go:154]
Reconciler: start to sync state

...

q

user@ubuntu:~/metrics-server$
```

The kubelet will restart and use the bootstrap configuration since you removed the kubelet.conf. The bootstrap configuration file location is declared in the kubelet service's `--bootstrap-kubeconfig` argument, and the permanent kubeconfig file is declared using the `--kubeconfig` argument.

Check if a new certificate signing request for the kubelet is present. This is the kubelet's request for a serving certificate:

```
user@ubuntu:~/metrics-server$ kubectl get csr
```

NAME	AGE	REQUESTOR	CONDITION
csr-2wnsr	2m6s	system:bootstrap:c2ihj5	Approved,Issued
csr-485fp	46s	system:node:ubuntu	Pending

```
user@ubuntu:~/metrics-server$
```

Use `kubectl certificate approve` on the Pending node certificate:

```
user@ubuntu:~/metrics-server$ kubectl certificate approve csr-485fp
```

certificatesigningrequest.certificates.k8s.io/csr-485fp approved

```
user@ubuntu:~/metrics-server$ kubectl get csr
```

NAME	AGE	REQUESTOR	CONDITION
csr-2wnsr	2m42s	system:bootstrap:c2ihj5	Approved,Issued
csr-485fp	82s	system:node:ubuntu	Approved,Issued

```
user@ubuntu:~/metrics-server$
```

The kubelet certificate is now running with a certificate signed by the cluster CA, meaning the API server and other services within the cluster will successfully validate any requests. This process needs to be performed for all kubelets.

3b. Reconfigure the metrics server to use the internal IP addresses of the kubelets

Another issue we need to address is how the Metrics Server communicates with the kubelets. The default behavior of the metrics server is to try and connect to the kubelets via their hostnames. Display the local node status and see how the node presents itself to the cluster:

```
user@ubuntu:~/metrics-server$ kubectl describe $(kubectl get node -o name) | grep -A2 Addresses
```

Addresses:

InternalIP: 192.168.229.132

Hostname: ubuntu

```
user@ubuntu:~/metrics-server$
```

Notice that the Node has a few IP address options and a hostname. Given that the metrics server runs in a Pod and may not have the ability to look up the hostname of the Kubelet's machine, so using the hostname (the metrics server default) is not ideal.

We can fix this problem by editing the Metrics Server deployment to pass the following argument to the Metrics Server:

- `--kubelet-preferred-address-types=InternalIP` - This will tell the Metrics Server to try to communicate using the Kubelet's InternalIP.

Display the deployments in the kube-system namespace:

```
user@ubuntu:~/metrics-server$ kubectl get deploy -n kube-system
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
coredns	2/2	2	2	24h
metrics-server	1/1	1	1	32m

```
user@ubuntu:~/metrics-server$
```

Make the changes to the Metrics Server deployment:

```
user@ubuntu:~/metrics-server$ kubectl edit -n kube-system deploy metrics-server  
...
```

In the editor add the `args` key and switches shown below:

```
...  
template:  
  metadata:  
    creationTimestamp: null  
    labels:  
      k8s-app: metrics-server  
    name: metrics-server  
  spec:  
    containers:  
      - args:  
        - --kubelet-preferred-address-types=InternalIP  
        - --cert-dir=/tmp  
        - --secure-port=4443  
        image: k8s.gcr.io/metrics-server-amd64:v0.3.6  
        imagePullPolicy: Always  
        name: metrics-server  
...  

```

Save your changes and exit:

```
...  
deployment.apps/metrics-server edited  
user@ubuntu:~/metrics-server$
```

Checking the pods again, you will see that the metrics server pod has been recreated and shows an age in seconds:

```
user@ubuntu:~/metrics-server$ kubectl get po -n kube-system  
  
NAME                                READY   STATUS    RESTARTS   AGE  
coredns-5644d7b6d9-gd8tv           1/1     Running   0           24h  
coredns-5644d7b6d9-p6rd6           1/1     Running   0           24h  
etcd-ubuntu                         1/1     Running   0           24h  
kube-apiserver-ubuntu               1/1     Running   0           24h  
kube-controller-manager-ubuntu      1/1     Running   0           24h  
kube-proxy-4c8tb                    1/1     Running   0           24h  
kube-scheduler-ubuntu               1/1     Running   0           24h  
metrics-server-c84799697-mhmqw      1/1     Running   0           36s  
weave-net-76twd                     2/2     Running   0           24h  
  
user@ubuntu:~/metrics-server$
```

It will take two or three minutes for the new Metrics Server to collect data but if you try a few times you should now see that top returns actual metrics. Use the Linux `watch` (not to be confused with `kubectl -watch`) tool to continually send a kubectl command. Use CTRL C once you see that the metrics are being reported:

```
user@ubuntu:~/metrics-server$ watch kubectl top node  
...
```

```
Every 2.0s: kubectl top node
```

```
Fri Jan 24 16:23:47 2020
```

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
ubuntu	140m	7%	1995Mi	52%

```
^C
```

```
user@ubuntu:~/metrics-server$
```

Try `kubectl top` on pods:

```
user@ubuntu:~/metrics-server$ kubectl top pod -n kube-system
```

NAME	CPU(cores)	MEMORY(bytes)
coredns-5644d7b6d9-gd8tv	2m	15Mi
coredns-5644d7b6d9-p6rd6	2m	24Mi
etcd-ubuntu	7m	59Mi
kube-apiserver-ubuntu	27m	301Mi
kube-controller-manager-ubuntu	6m	49Mi
kube-proxy-4c8tb	1m	22Mi
kube-scheduler-ubuntu	1m	21Mi
metrics-server-c84799697-mhmqw	1m	12Mi
weave-net-76twd	1m	94Mi

```
user@ubuntu:~/metrics-server$
```

Awesome, metrics up and running.

4. Autoscaling

Kubernetes 1.7 and later use the metrics-server to provide metrics used by the autoscaler controller. The base metrics used in most autoscaling scenarios are CPU and memory.

To test autoscaling we'll need:

- A target pod/container-image that we can load
- A pod we can use to drive the target pod
- A Horizontal Pod Autoscaler (HPA) to scale the target pod when it gets busy

To begin, create a simple php web server we can use as our target:

```
user@ubuntu:~/metrics-server$ mkdir ~/hpa && cd ~/hpa
```

```
user@ubuntu:~/hpa$ nano index.php && cat index.php
```

```
<?php
$x = 0.0001;
for ($i = 0; $i <= 1000000; $i++) {
    $x += sqrt($x);
}
echo "OK!";
?>
```

```
user@ubuntu:~/hpa$
```

This app will simply compute a million square roots each time it is hit.

Now create a Dockerfile to create a container image for the PHP server:

```
user@ubuntu:~/hpa$ nano Dockerfile && cat Dockerfile

FROM php:5-apache
ADD index.php /var/www/html/index.php
RUN chmod a+rx /var/www/html/index.php

user@ubuntu:~/hpa$
```

Build a container image for the PHP server:

```
user@ubuntu:~/hpa$ docker image build -t target .

Sending build context to Docker daemon 3.072kB
Step 1/3 : FROM php:5-apache
5-apache: Pulling from library/php
5e6ec7f28fb7: Pull complete
cf165947b5b7: Pull complete
7bd37682846d: Pull complete
99daf8e838e1: Pull complete
ae320713efba: Pull complete
ebcb99c48d8c: Pull complete
9867e71b4ab6: Pull complete
936eb418164a: Pull complete
bc298e7adaf7: Pull complete
ccd61b587bcd: Pull complete
b2d4b347f67c: Pull complete
56e9dde34152: Pull complete
9ad99b17eb78: Pull complete
Digest: sha256:0a40fd273961b99d8afe69a61a68c73c04bc0caa9de384d3b2dd9e7986eec86d
Status: Downloaded newer image for php:5-apache
---> 24c791995c1e
Step 2/3 : ADD index.php /var/www/html/index.php
---> 170b6d61f08f
Step 3/3 : RUN chmod a+rx /var/www/html/index.php
---> Running in 105f9f2ee939
Removing intermediate container 105f9f2ee939
---> ddafae5353b4
Successfully built ddafae5353b4
Successfully tagged target:latest

user@ubuntu:~/hpa$
```

N.B. Remember that Docker build places the PHP image on the node that you built it on! If you are running a multi-node Kubernetes cluster, make sure that any pods using this image run on a node where the image is present.

Now we can launch a standard single replica deployment with our new server:

```
user@ubuntu:~/hpa$ kubectl create deploy web1 --image=target --dry-run -o yaml > target-
deploy.yaml

user@ubuntu:~/hpa$ nano target-deploy.yaml

user@ubuntu:~/hpa$ cat target-deploy.yaml

apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: web1
    name: web1
spec:
  replicas: 1
  selector:
    matchLabels:
      app: web1
```

```

template:
  metadata:
    labels:
      app: web1
  spec:
    nodeName: ubuntu
    containers:
    - image: target
      imagePullPolicy: Never
      name: target
      resources:
        requests:
          cpu: "200m"

```

```

user@ubuntu:~/hpa$ kubectl create -f target-deploy.yaml

deployment.apps/web1 created

user@ubuntu:~/hpa$ kubectl expose deploy web1 --port 80

service/web1 exposed

user@ubuntu:~/hpa$

```

Our pod will request 200 mils of CPU (1/5 of a CPU) and we also set `imagePullPolicy` to "Never" to prevent the kubelet from asking Docker to download the image (which it would not find since we have just built it locally).

N.B. A `nodeName` key keeps this deployment's pods from running on a node that doesn't have the PHP image you built. Make sure to set it to the hostname of the node that has the PHP image.

List your resources:

```

user@ubuntu:~/hpa$ kubectl get deploy,rs,po,svc

```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/web1	1/1	1	1	47s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/web1-75c54b6579	1	1	1	47s

NAME	READY	STATUS	RESTARTS	AGE
pod/web1-75c54b6579-lqf7z	1/1	Running	0	47s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	24h
service/web1	ClusterIP	10.103.152.45	<none>	80/TCP	44s

```

user@ubuntu:~/hpa$

```

So we now have our deployment with a replica set, one pod and a service (created because of the `--expose` switch).

Now let's create an HPA for our deployment:

```

user@ubuntu:~/hpa$ kubectl autoscale deployment web1 --cpu-percent=50 --min=1 --max=5

horizontalpodautoscaler.autoscaling/web1 autoscaled

user@ubuntu:~/hpa$

```

Our HPA will now scale up when our set of pods exceed 50% of their desired CPU resources, in this case 100 mils. Right now we have one pod with a request for 20% of a CPU, or 200 mils. Display the HPA. You may see the Targets for the pod read for a few minutes until the autoscaler communicates with the metrics server:

```

user@ubuntu:~/hpa$ kubectl get hpa

```


NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
web1	Deployment/web1	0%/50%	1	5	1	63s

```
user@ubuntu:~/hpa$
```

Now display the resources used by your pod:

```
user@ubuntu:~/hpa$ kubectl top pod
```

NAME	CPU(cores)	MEMORY(bytes)
web1-75c54b6579-lqf7z	1m	8Mi

```
user@ubuntu:~/hpa$
```

Our pod is using 1 mil of its 200 mils. It will need to get to 100 mils before the HPA will add another pod to the deployment. Let's create some load!

Run an interactive busybox pod to hit the PHP server from:

```
user@ubuntu:~/hpa$ kubectl run --generator=run-pod/v1 -it driver --image=busybox
```

If you don't see a command prompt, try pressing enter.

```
/ #
```

Now start a loop that will make continuous requests of our 1mm square root service:

```
/ # while true; do wget -q -O- http://web1.default.svc.cluster.local; done
OK!
OK!
OK!

...
```

While the requests are running, **open a new terminal** and check the pod utilization with top:

```
user@ubuntu:~$ kubectl top pod
```

NAME	CPU(cores)	MEMORY(bytes)
web1-75c54b6579-lqf7z	1m	8Mi

```
user@ubuntu:~$ kubectl top pod
```

NAME	CPU(cores)	MEMORY(bytes)
driver	7m	1Mi
web1-75c54b6579-lqf7z	928m	10Mi

```
user@ubuntu:~$
```

Now over 900 mils, that is well over the threshold of 100 mils, why don't we have more pods reporting metrics? The HPA will not respond to spikes because this would trash the cluster, creating and deleting pods wastefully. Rather the HPA waits for a configurable period (defined by Controller Manager settings, defaulting to about 3 minutes) and then begins to scale.

Run top several more times until you see scaling activity, display the HPA status now and then also:

```
user@ubuntu:~$ kubectl top pod
```

NAME	CPU(cores)	MEMORY(bytes)
driver	7m	1Mi

```
web1-75c54b6579-lqf7z    928m    10Mi

user@ubuntu:~$ kubectl get hpa

NAME      REFERENCE          TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
web1      Deployment/web1    83%/50%   1         5         5         3m11s

user@ubuntu:~$ kubectl top pod

NAME                CPU(cores)   MEMORY(bytes)
driver              10m          2Mi
web1-75c54b6579-gjss2 170m         10Mi
web1-75c54b6579-k5tqp 164m         10Mi
web1-75c54b6579-lqf7z 203m         10Mi
web1-75c54b6579-psm69 130m         10Mi
web1-75c54b6579-w6d4c 170m         10Mi

user@ubuntu:~$
```

You should see the HPA scale the pods out to the limit of 5 configured. Once those pods communicate with the metrics server, the start showing up on `kubectl top` output. Display your deployment:

```
user@ubuntu:~$ kubectl get deploy

NAME    READY   UP-TO-DATE   AVAILABLE   AGE
web1    5/5     5            5           4m32s

user@ubuntu:~$
```

Note that the only thing the HPA actually does is change the replica count on the deployment. The deployment and replica set do the rest.

Use `control C` to stop the driver loop and exit the pod:

```
...

OK!
OK!

^C

/ # exit

Session ended, resume using 'kubectl attach driver -c driver -i -t' command when the pod is
running

user@ubuntu:~/hpa$
```

Examine the metrics and HPA:

```
user@ubuntu:~/hpa$ kubectl get hpa

NAME      REFERENCE          TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
web1      Deployment/web1    90%/50%   1         5         5         3m58s

user@ubuntu:~/hpa$ kubectl top pod

NAME                CPU(cores)   MEMORY(bytes)
driver              10m          2Mi
web1-75c54b6579-gjss2 122m         10Mi
web1-75c54b6579-k5tqp 182m         10Mi
web1-75c54b6579-lqf7z 202m         10Mi
web1-75c54b6579-psm69 211m         10Mi
web1-75c54b6579-w6d4c 192m         10Mi
```

```
user@ubuntu:~/hpa$
```

The HPA will wait to ensure traffic does not return for a few minutes before scaling down. While you wait for the HPA to scale down, describe the resource:

```
user@ubuntu:~/hpa$ kubectl describe hpa web1
```

```
Name: web1
Namespace: default
Labels: <none>
Annotations: <none>
CreationTimestamp: Fri, 24 Jan 2020 16:28:17 -0800
Reference: Deployment/web1
Metrics: ( current / target )
  resource cpu on pods  (as a percentage of request): 0% (1m) / 50%
Min replicas: 1
Max replicas: 5
Deployment pods: 5 current / 5 desired
Conditions:
  Type           Status  Reason
  ----           -
  AbleToScale    True    ScaleDownStabilized recent recommendations were higher than current
  one, applying the highest recent recommendation
  ScalingActive  True    ValidMetricFound the HPA was able to successfully calculate a
  replica count from cpu resource utilization (percentage of request)
  ScalingLimited True    TooManyReplicas the desired replica count is more than the
  maximum replica count
Events:
  Type           Reason              Age           From                               Message
  ----           -
  Warning        FailedGetResourceMetric 4m28s        horizontal-pod-autoscaler    did not receive
  metrics for any ready pods
  Warning        FailedComputeMetricsReplicas 4m28s        horizontal-pod-autoscaler    invalid metrics (1
  invalid out of 1), first error is: failed to get cpu utilization: did not receive metrics for
  any ready pods
  Normal        SuccessfulRescale      3m12s        horizontal-pod-autoscaler    New size: 4; reason:
  cpu resource utilization (percentage of request) above target
  Normal        SuccessfulRescale      2m57s        horizontal-pod-autoscaler    New size: 5; reason:
  cpu resource utilization (percentage of request) above target

user@ubuntu:~/hpa$
```

Notice the scaling events are reported at the bottom of the description.

The HPA abides by a setting called `horizontal-pod-autoscaler-downscale-stabilization` which is passed as a flag to the cluster's kube-controller-manager. This argument prevents an HPA from scaling down for 5 minutes by default. After 5 minutes, the HPA will scale the deployment back down to one pod.

Use `watch` to repeatedly run the `kubectl get hpa` and `kubectl top pod` commands:

```
user@ubuntu:~/hpa$ watch 'kubectl get hpa && kubectl top pod'
```

```
Every 2.0s: kubectl get hpa && kubectl top pod          Fri Jan 24 16:33:44 2020
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
web1	Deployment/web1	0%/50%	1	5	5	5m27s
NAME		CPU(cores)	MEMORY(bytes)			
driver		0m	1Mi			
web1-75c54b6579-gjss2		1m	10Mi			
web1-75c54b6579-k5tqp		1m	10Mi			
web1-75c54b6579-lqf7z		1m	10Mi			
web1-75c54b6579-psm69		1m	10Mi			
web1-75c54b6579-w6d4c		1m	10Mi			

```
user@ubuntu:~/hpa$
```

After about five minutes, you will see the deployment scale down. Use `Ctrl c` to exit the watch when you see the scale down event:

```
Every 2.0s: kubectl get hpa && kubectl top pod          Fri Jan 24 16:37:47 2020

NAME      REFERENCE          TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
web1      Deployment/web1    0%/50%   1         5         5          9m30s
NAME      CPU(cores)  MEMORY(bytes)
driver    0m          1Mi
web1-75c54b6579-lqf7z  1m          10Mi

^C
user@ubuntu:~/hpa$
```

Describe the HPA once more and you will see a scaling event at the bottom that reports a new size and the reason:

```
user@ubuntu:~/hpa$ kubectl describe hpa | grep Events -A10

Events:
  Type     Reason                                     Age    From                                     Message
  ----     -
  Warning   FailedGetResourceMetric                  9m50s  horizontal-pod-autoscaler              did not receive
metrics for any ready pods
  Warning   FailedComputeMetricsReplicas            9m50s  horizontal-pod-autoscaler              invalid metrics (1
invalid out of 1), first error is: failed to get cpu utilization: did not receive metrics for
any ready pods
  Normal    SuccessfulRescale                        8m34s  horizontal-pod-autoscaler              New size: 4; reason:
cpu resource utilization (percentage of request) above target
  Normal    SuccessfulRescale                        8m19s  horizontal-pod-autoscaler              New size: 5; reason:
cpu resource utilization (percentage of request) above target
  Normal    SuccessfulRescale                        46s    horizontal-pod-autoscaler              New size: 1; reason:
All metrics below target

user@ubuntu:~/hpa$
```

Your cluster is now capable of dynamically sizing its elements based on metrics measured from incoming workloads.

5. Clean up

List all of your resources:

```
user@ubuntu:~/hpa$ kubectl get all -o=name

pod/driver
pod/web1-75c54b6579-lqf7z
service/kubernetes
service/web1
deployment.apps/web1
replicaset.apps/web1-75c54b6579
horizontalpodautoscaler.autoscaling/web1

user@ubuntu:~/hpa$
```

Now delete everything you created for this lab (but don't delete the "kubernetes" service!!):

```
user@ubuntu:~/hpa$ kubectl delete service/web1 deploy/web1 pod/driver hpa/web1

service "web1" deleted
deployment.apps "web1" deleted
pod "driver" deleted
horizontalpodautoscaler.autoscaling "web1" deleted
```

```
user@ubuntu:~/hpa$
```

Congratulations, you have completed the lab!

Copyright (c) 2015-2020 RX-M LLC, Cloud Native Consulting, all rights reserved